

# 20104028

## Heamnath N

### Importing Libraries

In [1]:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

### Importing Datasets

In [2]:

```
df=pd.read_csv("madrid_2002.csv")
df
```

Out[2]:

	<b>date</b>	<b>BEN</b>	<b>CO</b>	<b>EBE</b>	<b>MXY</b>	<b>NMHC</b>	<b>NO_2</b>	<b>NOx</b>	<b>OXY</b>	<b>O_3</b>	<b>PM10</b>	<b>PXY</b>
<b>0</b>	2002-04-01 01:00:00	NaN	1.39	NaN	NaN	NaN	145.100006	352.100006	NaN	6.54	41.990002	NaN
<b>1</b>	2002-04-01 01:00:00	1.93	0.71	2.33	6.20	0.15	98.150002	153.399994	2.67	6.85	20.980000	2.53
<b>2</b>	2002-04-01 01:00:00	NaN	0.80	NaN	NaN	NaN	103.699997	134.000000	NaN	13.01	28.440001	NaN
<b>3</b>	2002-04-01 01:00:00	NaN	1.61	NaN	NaN	NaN	97.599998	268.000000	NaN	5.12	42.180000	NaN
<b>4</b>	2002-04-01 01:00:00	NaN	1.90	NaN	NaN	NaN	92.089996	237.199997	NaN	7.28	76.330002	NaN
...	...	...	...	...	...	...	...	...	...	...	...	...
<b>217291</b>	2002-11-01 00:00:00	4.16	1.14	NaN	NaN	NaN	81.080002	265.700012	NaN	7.21	36.750000	NaN
<b>217292</b>	2002-11-01 00:00:00	3.67	1.73	2.89	NaN	0.38	113.900002	373.100006	NaN	5.66	63.389999	NaN
<b>217293</b>	2002-11-01 00:00:00	1.37	0.58	1.17	2.37	0.15	65.389999	107.699997	1.30	9.11	9.640000	0.94

	date	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3	PM10	PXY
217294	2002-11-01 00:00:00	4.51	0.91	4.83	10.99	NaN	149.800003	202.199997	1.00	5.75	NaN	5.52
217295	2002-11-01 00:00:00	3.11	1.17	3.00	7.77	0.26	80.110001	180.300003	2.25	7.38	29.240000	3.35

217296 rows × 16 columns

## Data Cleaning and Data Preprocessing

In [3]: `df=df.dropna()`In [4]: `df.columns`Out[4]: `Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3', 'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'], dtype='object')`In [5]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 32381 entries, 1 to 217295
Data columns (total 16 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   date        32381 non-null   object 
 1   BEN          32381 non-null   float64
 2   CO           32381 non-null   float64
 3   EBE          32381 non-null   float64
 4   MXY          32381 non-null   float64
 5   NMHC         32381 non-null   float64
 6   NO_2         32381 non-null   float64
 7   NOx          32381 non-null   float64
 8   OXY          32381 non-null   float64
 9   O_3           32381 non-null   float64
 10  PM10         32381 non-null   float64
 11  PXY          32381 non-null   float64
 12  SO_2         32381 non-null   float64
 13  TCH          32381 non-null   float64
 14  TOL          32381 non-null   float64
 15  station      32381 non-null   int64  
dtypes: float64(14), int64(1), object(1)
memory usage: 4.2+ MB
```

In [6]: `data=df[['CO', 'station']]`  
`data`Out[6]: 

	CO	station
1	0.71	28079035

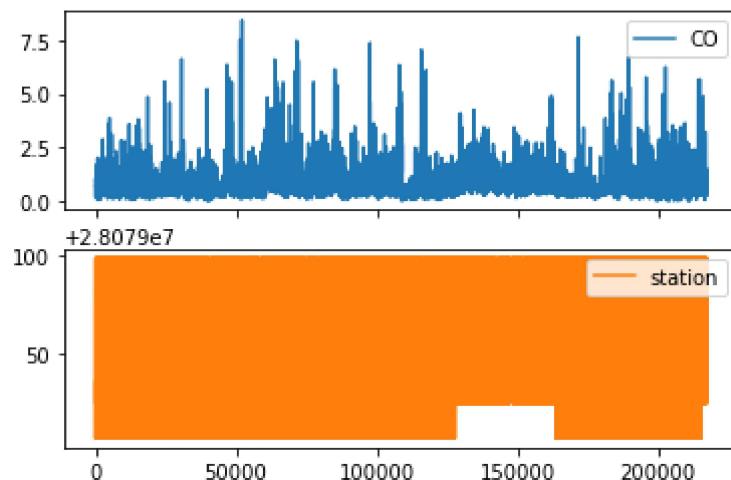
	CO	station
5	0.72	28079006
22	0.80	28079024
24	1.04	28079099
26	0.53	28079035
...	...	...
217269	0.28	28079024
217271	1.30	28079099
217273	0.97	28079035
217293	0.58	28079024
217295	1.17	28079099

32381 rows × 2 columns

## Line chart

In [7]: `data.plot.line(subplots=True)`

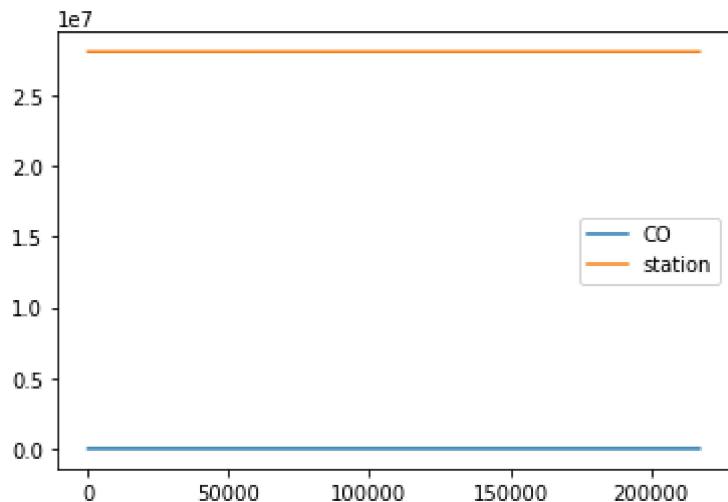
Out[7]: `array([<AxesSubplot:>, <AxesSubplot:>], dtype=object)`



## Line chart

In [8]: `data.plot.line()`

Out[8]: `<AxesSubplot:>`

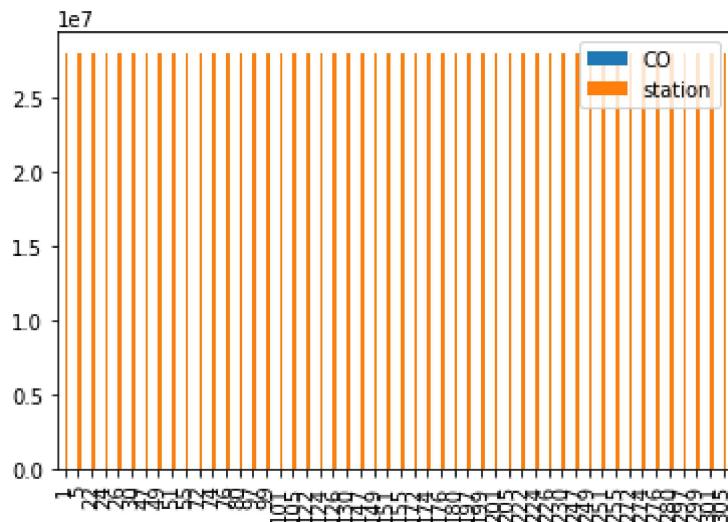


## Bar chart

```
In [9]: b=data[0:50]
```

```
In [10]: b.plot.bar()
```

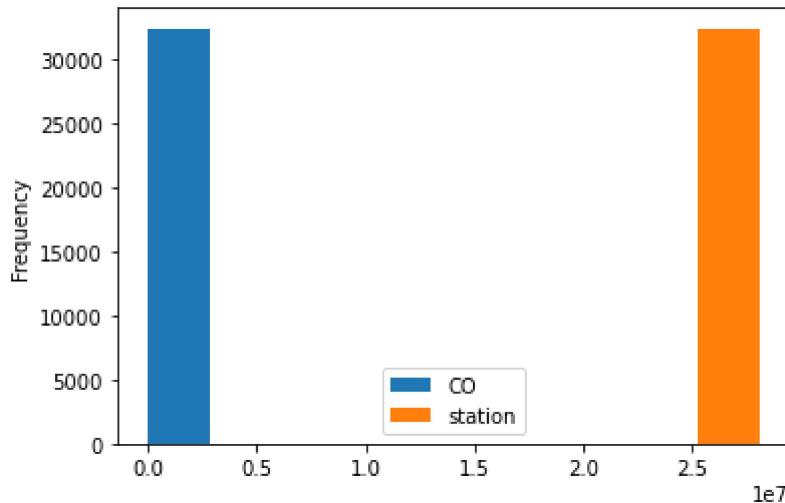
```
Out[10]: <AxesSubplot:>
```



## Histogram

```
In [11]: data.plot.hist()
```

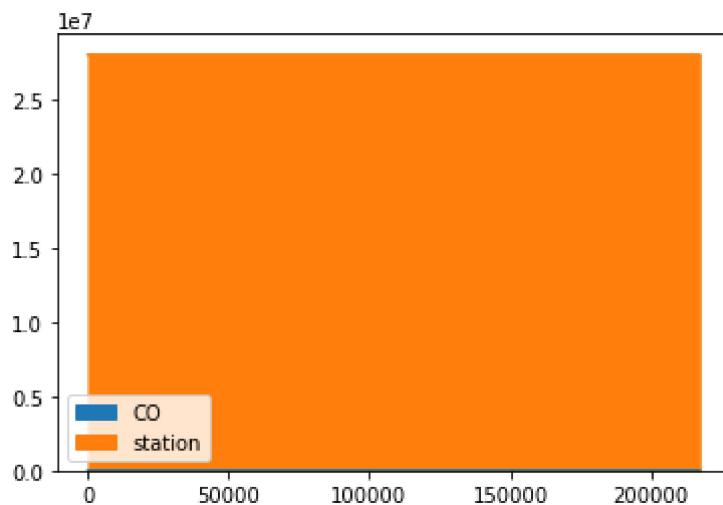
```
Out[11]: <AxesSubplot:ylabel='Frequency'>
```



## Area chart

In [12]: `data.plot.area()`

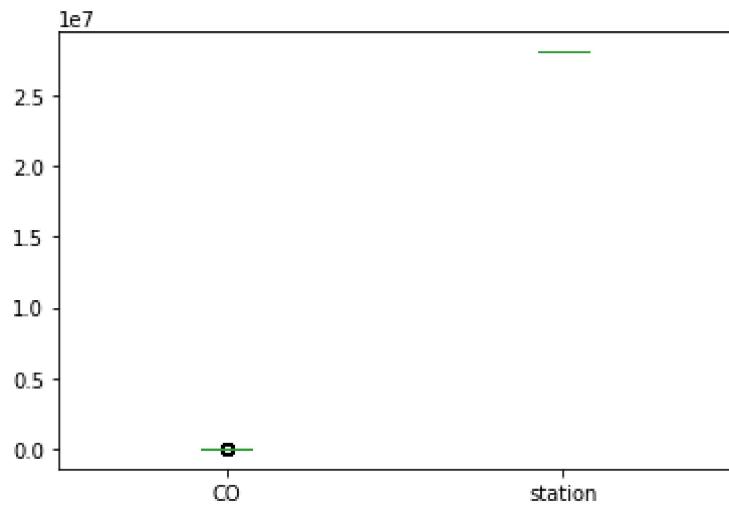
Out[12]: <AxesSubplot:>



## Box chart

In [13]: `data.plot.box()`

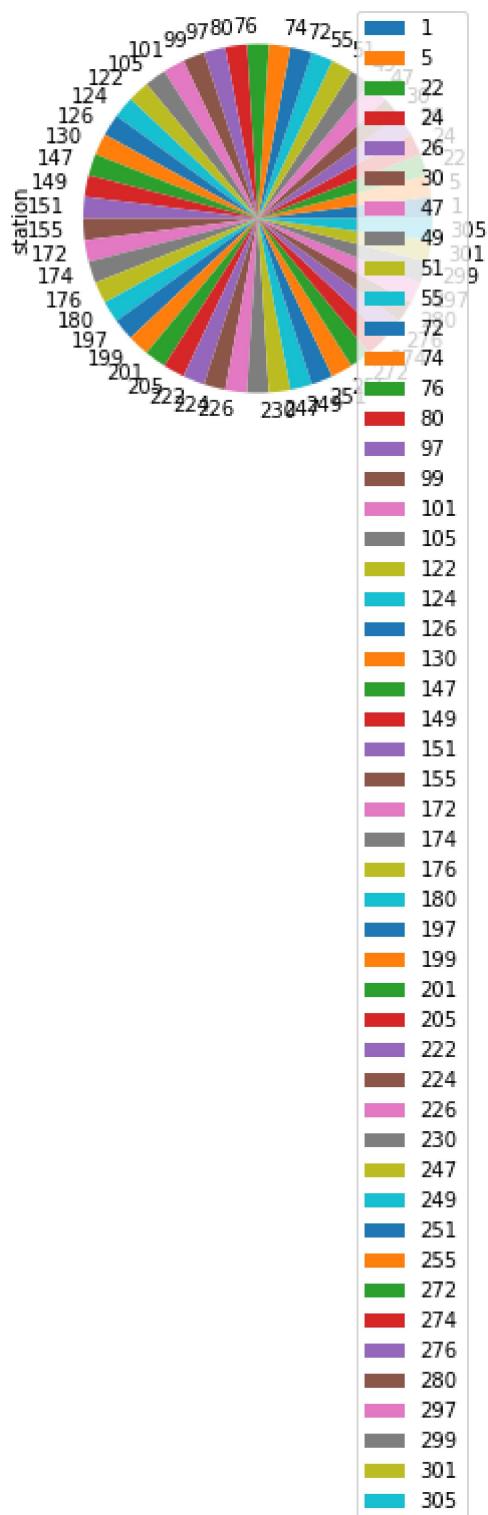
Out[13]: <AxesSubplot:>



## Pie chart

```
In [14]: b.plot.pie(y='station' )
```

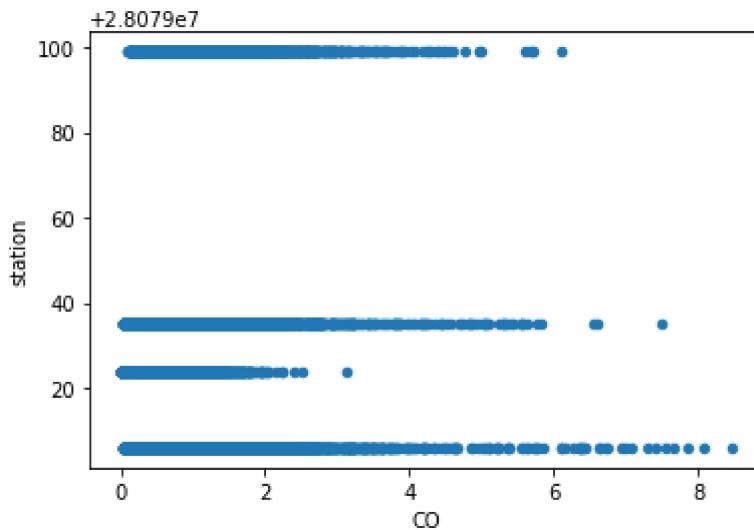
```
Out[14]: <AxesSubplot:ylabel='station'>
```



## Scatter chart

In [15]: `data.plot.scatter(x='CO' ,y='station')`

Out[15]: <AxesSubplot:xlabel='CO', ylabel='station'>



In [16]:

`df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 32381 entries, 1 to 217295
Data columns (total 16 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   date        32381 non-null   object 
 1   BEN          32381 non-null   float64
 2   CO           32381 non-null   float64
 3   EBE          32381 non-null   float64
 4   MXY          32381 non-null   float64
 5   NMHC         32381 non-null   float64
 6   NO_2          32381 non-null   float64
 7   NOX          32381 non-null   float64
 8   OXY          32381 non-null   float64
 9   O_3           32381 non-null   float64
 10  PM10         32381 non-null   float64
 11  PXY          32381 non-null   float64
 12  SO_2          32381 non-null   float64
 13  TCH          32381 non-null   float64
 14  TOL          32381 non-null   float64
 15  station       32381 non-null   int64  
dtypes: float64(14), int64(1), object(1)
memory usage: 4.2+ MB
```

In [17]:

`df.describe()`

Out[17]:

	BEN	CO	EBE	MXY	NMHC	NO_2	NO
<b>count</b>	32381.000000	32381.000000	32381.000000	32381.000000	32381.000000	32381.000000	32381.000000
<b>mean</b>	2.479155	0.787323	2.914004	7.013636	0.155827	58.936796	126.009340
<b>std</b>	2.280959	0.610810	2.667881	6.774365	0.135731	31.472733	114.035071
<b>min</b>	0.180000	0.000000	0.180000	0.190000	0.000000	0.890000	1.710000
<b>25%</b>	0.970000	0.420000	1.140000	2.420000	0.080000	35.660000	48.720000
<b>50%</b>	1.840000	0.620000	2.130000	5.140000	0.130000	57.160000	96.830000
<b>75%</b>	3.250000	0.980000	3.830000	9.420000	0.200000	78.769997	167.500000

	BEN	CO	EBE	MXY	NMHC	NO_2	NOx
max	32.660000	8.460000	41.740002	99.879997	2.700000	263.600006	1336.000000

In [18]:

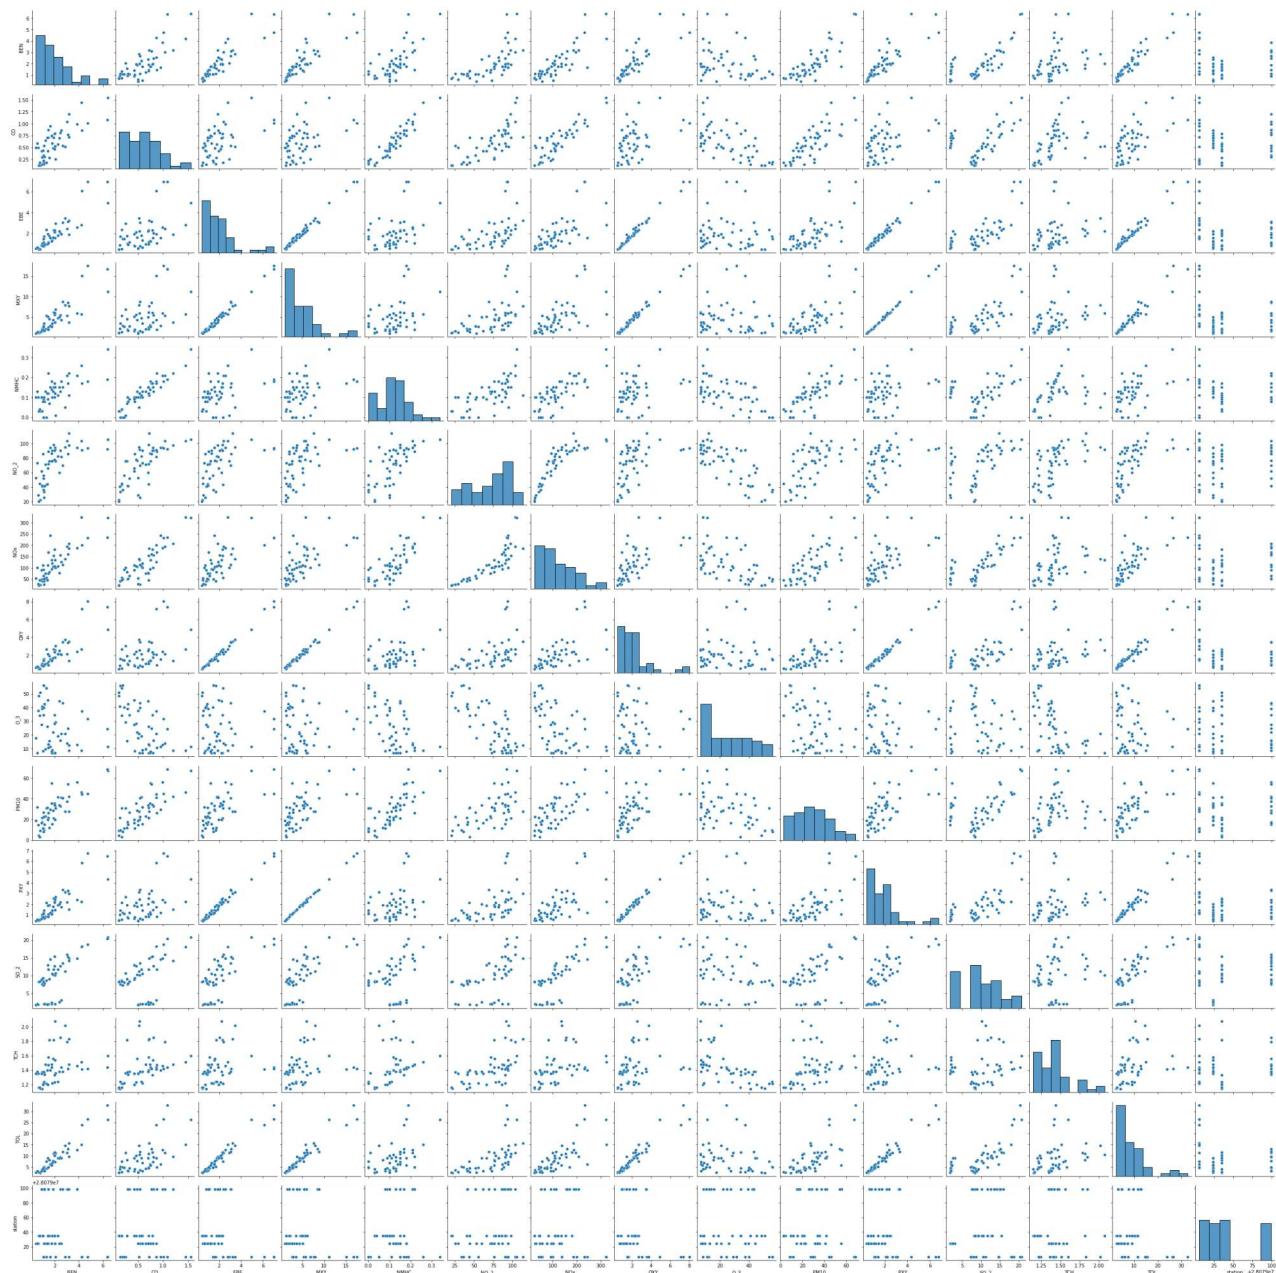
```
df1=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
       'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]
```

## EDA AND VISUALIZATION

In [19]:

```
sns.pairplot(df1[0:50])
```

Out[19]: &lt;seaborn.axisgrid.PairGrid at 0x1d32b1e36d0&gt;

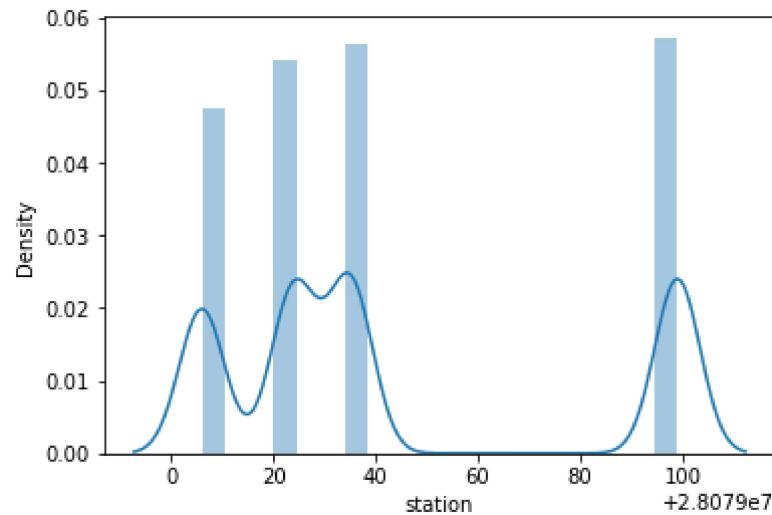


In [20]:

```
sns.distplot(df1['station'])
```

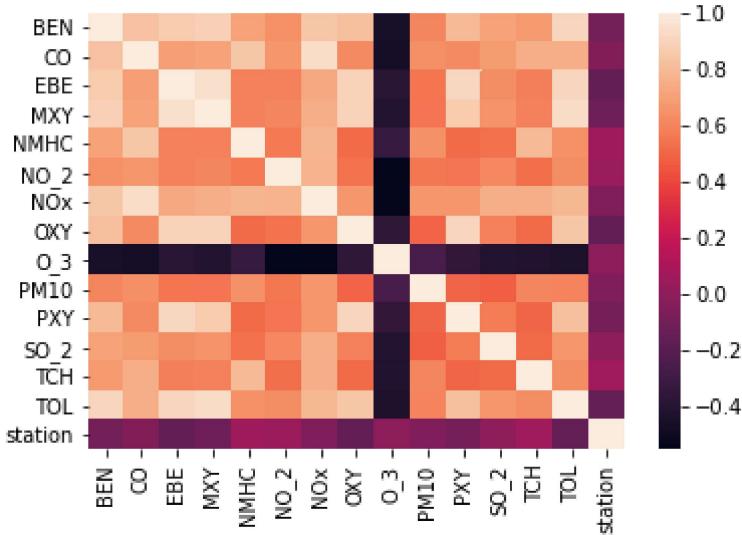
```
C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning:
`distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
    warnings.warn(msg, FutureWarning)
```

Out[20]: <AxesSubplot:xlabel='station', ylabel='Density'>



In [21]: `sns.heatmap(df1.corr())`

Out[21]: <AxesSubplot:>



## TO TRAIN THE MODEL AND MODEL BUILDING

In [22]:

```
x=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
       'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
y=df['station']
```

```
In [23]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

# Linear Regression

```
In [24]: from sklearn.linear_model import LinearRegression  
lr=LinearRegression()  
lr.fit(x_train,y_train)
```

```
Out[24]: LinearRegression()
```

```
In [25]: lr.intercept_
```

```
Out[25]: 28078989.001570065
```

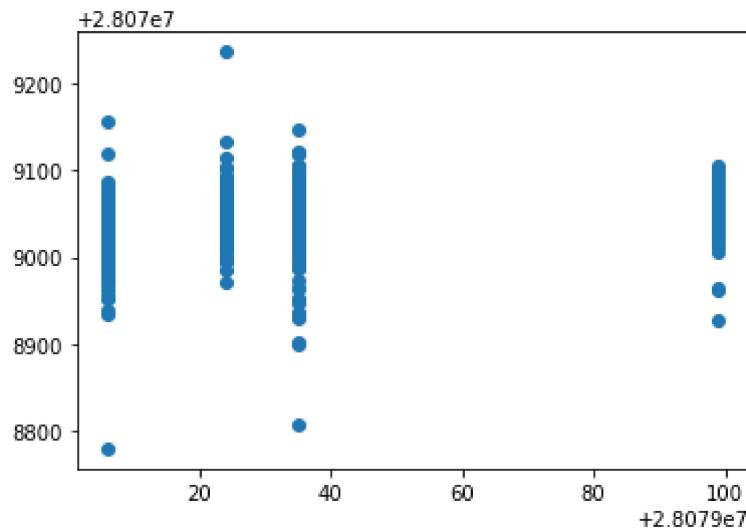
```
In [26]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
coeff
```

```
Out[26]:
```

	Co-efficient
BEN	1.682400
CO	-13.299938
EBE	-11.699361
MXY	3.944189
NMHC	83.582973
NO_2	0.247211
NOx	-0.097943
OXY	-4.829133
O_3	-0.030373
PM10	-0.132962
PXY	7.568084
SO_2	0.582677
TCH	42.944334
TOL	-1.307853

```
In [27]: prediction =lr.predict(x_test)  
plt.scatter(y_test,prediction)
```

```
Out[27]: <matplotlib.collections.PathCollection at 0x1d33a9ea6a0>
```



## ACCURACY

```
In [28]: lr.score(x_test,y_test)
```

```
Out[28]: 0.20446058578801884
```

```
In [29]: lr.score(x_train,y_train)
```

```
Out[29]: 0.19625702297940784
```

## Ridge and Lasso

```
In [30]: from sklearn.linear_model import Ridge,Lasso
```

```
In [31]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

```
Out[31]: Ridge(alpha=10)
```

## Accuracy(Ridge)

```
In [32]: rr.score(x_test,y_test)
```

```
Out[32]: 0.20400914542939697
```

```
In [33]: rr.score(x_train,y_train)
```

```
Out[33]: 0.1960481691260696
```

```
In [34]: la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

Out[34]: Lasso(alpha=10)

```
In [35]: la.score(x_train,y_train)
```

Out[35]: 0.056416235247383595

## Accuracy(Lasso)

```
In [36]: la.score(x_test,y_test)
```

Out[36]: 0.05954671177337789

```
In [37]: from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

Out[37]: ElasticNet()

```
In [38]: en.coef_
```

Out[38]: array([ 8.33976815e-01, 0.00000000e+00, -2.87847208e+00, 1.45141342e+00,
 2.24380188e-01, 2.21290459e-01, -2.83345497e-02, -2.35327031e+00,
 -2.86286966e-02, 1.60240838e-03, 2.25256815e+00, 3.85747481e-01,
 1.09063241e+00, -1.08607526e+00])

```
In [39]: en.intercept_
```

Out[39]: 28079038.96998315

```
In [40]: prediction=en.predict(x_test)
```

```
In [41]: en.score(x_test,y_test)
```

Out[41]: 0.10310757444971508

## Evaluation Metrics

```
In [42]: from sklearn import metrics
print(metrics.mean_absolute_error(y_test,prediction))
print(metrics.mean_squared_error(y_test,prediction))
print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
28.596830616620665  
1120.8289730725503  
33.47878392463726
```

# Logistic Regression

```
In [43]: from sklearn.linear_model import LogisticRegression
```

```
In [44]: feature_matrix=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOX', 'OXY', 'O_3',  
'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]  
target_vector=df[ 'station']
```

```
In [45]: feature_matrix.shape
```

```
Out[45]: (32381, 14)
```

```
In [46]: target_vector.shape
```

```
Out[46]: (32381,)
```

```
In [47]: from sklearn.preprocessing import StandardScaler
```

```
In [48]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [49]: logr=LogisticRegression(max_iter=10000)  
logr.fit(fs,target_vector)
```

```
Out[49]: LogisticRegression(max_iter=10000)
```

```
In [50]: observation=[[1,2,3,4,5,6,7,8,9,10,11,12,13,14]]
```

```
In [51]: prediction=logr.predict(observation)  
print(prediction)
```

```
[28079035]
```

```
In [52]: logr.classes_
```

```
Out[52]: array([28079006, 28079024, 28079035, 28079099], dtype=int64)
```

```
In [53]: logr.score(fs,target_vector)
```

```
Out[53]: 0.8480899292795158
```

```
In [54]: logr.predict_proba(observation)[0][0]
```

```
Out[54]: 2.5638972732451705e-10
```

```
In [55]: logr.predict_proba(observation)
```

```
Out[55]: array([[2.56389727e-10, 3.44199742e-71, 1.00000000e+00, 1.43898646e-13]])
```

## Random Forest

```
In [56]: from sklearn.ensemble import RandomForestClassifier
```

```
In [57]: rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[57]: RandomForestClassifier()
```

```
In [58]: parameters={'max_depth':[1,2,3,4,5],
                 'min_samples_leaf':[5,10,15,20,25],
                 'n_estimators':[10,20,30,40,50]
                }
```

```
In [59]: from sklearn.model_selection import GridSearchCV
grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")
grid_search.fit(x_train,y_train)
```

```
Out[59]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
                      param_grid={'max_depth': [1, 2, 3, 4, 5],
                                  'min_samples_leaf': [5, 10, 15, 20, 25],
                                  'n_estimators': [10, 20, 30, 40, 50]},
                      scoring='accuracy')
```

```
In [60]: grid_search.best_score_
```

```
Out[60]: 0.7747286684902497
```

```
In [61]: rfc_best=grid_search.best_estimator_
```

```
In [62]: from sklearn.tree import plot_tree
```

```
plt.figure(figsize=(80,40))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['a','b','c','d'])
```

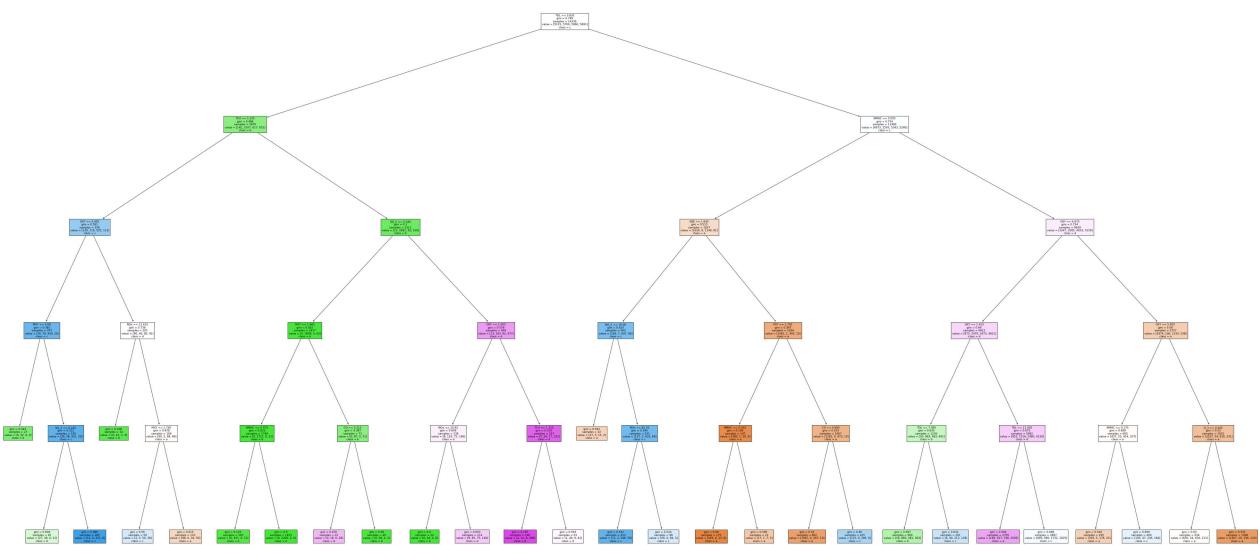
```
Out[62]: [Text(2010.9056603773586, 1993.2, 'TOL <= 3.835\nngini = 0.749\nnsamples = 14336\nvalue = [5115, 5700, 5960, 5891]\nclass = c'),
          Text(884.377358490566, 1630.8000000000002, 'TCH <= 1.225\nngini = 0.486\nnsamples = 2870\nvalue = [142, 3107, 617, 651]\nclass = b')]
```

```

Text(336.9056603773585, 1268.4, 'OXY <= 0.905\ngini = 0.591\nsamples = 558\nvalue = [12
9, 116, 525, 112]\nclass = c'),
Text(168.45283018867926, 906.0, 'MXY <= 0.69\ngini = 0.381\nsamples = 357\nvalue = [39,
70, 440, 20]\nclass = c'),
Text(84.22641509433963, 543.5999999999999, 'gini = 0.343\nsamples = 27\nvalue = [0, 32,
9, 0]\nclass = b'),
Text(252.67924528301887, 543.5999999999999, 'SO_2 <= 6.165\ngini = 0.322\nsamples = 330
\nvalue = [39, 38, 431, 20]\nclass = c'),
Text(168.45283018867926, 181.1999999999982, 'gini = 0.609\nsamples = 45\nvalue = [27,
38, 0, 12]\nclass = b'),
Text(336.9056603773585, 181.1999999999982, 'gini = 0.086\nsamples = 285\nvalue = [12,
0, 431, 8]\nclass = c'),
Text(505.35849056603774, 906.0, 'NOx <= 13.615\ngini = 0.736\nsamples = 201\nvalue = [9
0, 46, 85, 92]\nclass = d'),
Text(421.13207547169816, 543.5999999999999, 'gini = 0.198\nsamples = 32\nvalue = [0, 4
1, 1, 4]\nclass = b'),
Text(589.5849056603774, 543.5999999999999, 'MXY <= 1.745\ngini = 0.678\nsamples = 169\n
value = [90, 5, 84, 88]\nclass = a'),
Text(505.35849056603774, 181.1999999999982, 'gini = 0.56\nsamples = 59\nvalue = [2, 5,
50, 38]\nclass = c'),
Text(673.811320754717, 181.1999999999982, 'gini = 0.615\nsamples = 110\nvalue = [88,
0, 34, 50]\nclass = a'),
Text(1431.8490566037738, 1268.4, 'SO_2 <= 5.345\ngini = 0.3\nsamples = 2312\nvalue = [1
3, 2991, 92, 539]\nclass = b'),
Text(1094.9433962264152, 906.0, 'MXY <= 1.865\ngini = 0.042\nsamples = 1817\nvalue =
[0, 2808, 0, 62]\nclass = b'),
Text(926.4905660377359, 543.5999999999999, 'NMHC <= 0.075\ngini = 0.022\nsamples = 1746
\nvalue = [0, 2721, 0, 31]\nclass = b'),
Text(842.2641509433963, 181.1999999999982, 'gini = 0.125\nsamples = 293\nvalue = [0, 4
31, 0, 31]\nclass = b'),
Text(1010.7169811320755, 181.1999999999982, 'gini = 0.0\nsamples = 1453\nvalue = [0, 2
290, 0, 0]\nclass = b'),
Text(1263.3962264150944, 543.5999999999999, 'CO <= 0.315\ngini = 0.387\nsamples = 71\nv
alue = [0, 87, 0, 31]\nclass = b'),
Text(1179.169811320755, 181.1999999999982, 'gini = 0.476\nsamples = 22\nvalue = [0, 1
8, 0, 28]\nclass = d'),
Text(1347.622641509434, 181.1999999999982, 'gini = 0.08\nsamples = 49\nvalue = [0, 69,
0, 3]\nclass = b'),
Text(1768.754716981132, 906.0, 'OXY <= 1.005\ngini = 0.539\nsamples = 495\nvalue = [13,
183, 92, 477]\nclass = d'),
Text(1600.301886792453, 543.5999999999999, 'NOx <= 15.92\ngini = 0.659\nsamples = 238\nv
alue = [8, 129, 75, 146]\nclass = d'),
Text(1516.0754716981132, 181.1999999999982, 'gini = 0.0\nsamples = 24\nvalue = [0, 44,
0, 0]\nclass = b'),
Text(1684.5283018867926, 181.1999999999982, 'gini = 0.653\nsamples = 214\nvalue = [8,
85, 75, 146]\nclass = d'),
Text(1937.2075471698115, 543.5999999999999, 'TCH <= 1.325\ngini = 0.319\nsamples = 257
\nvalue = [5, 54, 17, 331]\nclass = d'),
Text(1852.9811320754718, 181.1999999999982, 'gini = 0.155\nsamples = 196\nvalue = [4,
14, 8, 289]\nclass = d'),
Text(2021.433962264151, 181.1999999999982, 'gini = 0.593\nsamples = 61\nvalue = [1, 4
0, 9, 42]\nclass = d'),
Text(3137.433962264151, 1630.8000000000002, 'NMHC <= 0.055\ngini = 0.734\nsamples = 114
66\nvalue = [4973, 2593, 5343, 5240]\nclass = c'),
Text(2484.679245283019, 1268.4, 'EBE <= 1.645\ngini = 0.515\nsamples = 1827\nvalue = [1
626, 8, 1140, 81]\nclass = a'),
Text(2189.8867924528304, 906.0, 'NO_2 <= 19.06\ngini = 0.424\nsamples = 563\nvalue = [1
64, 7, 647, 66]\nclass = c'),
Text(2105.6603773584907, 543.5999999999999, 'gini = 0.591\nsamples = 32\nvalue = [27,
5, 15, 2]\nclass = a'),
Text(2274.11320754717, 543.5999999999999, 'NOx <= 82.16\ngini = 0.394\nsamples = 531\nv
alue = [137, 2, 632, 64]\nclass = c'),
Text(2189.8867924528304, 181.1999999999982, 'gini = 0.332\nsamples = 433\nvalue = [71,
2, 548, 59]\nclass = c'),
Text(2358.33962264151, 181.1999999999982, 'gini = 0.524\nsamples = 98\nvalue = [66, 0,

```

```
84, 5]\nclass = c'),  
    Text(2779.471698113208, 906.0, 'PXY <= 1.745\ngini = 0.387\nsamples = 1264\nvalue = [14  
62, 1, 493, 15]\nclass = a'),  
    Text(2611.0188679245284, 543.5999999999999, 'NMHC <= 0.035\ngini = 0.158\nsamples = 197  
\nvalue = [280, 1, 20, 5]\nclass = a'),  
    Text(2526.7924528301887, 181.1999999999982, 'gini = 0.09\nsamples = 175\nvalue = [263,  
0, 13, 0]\nclass = a'),  
    Text(2695.245283018868, 181.1999999999982, 'gini = 0.596\nsamples = 22\nvalue = [17,  
1, 7, 5]\nclass = a'),  
    Text(2947.924528301887, 543.5999999999999, 'CO <= 0.665\ngini = 0.415\nsamples = 1067\n  
value = [1182, 0, 473, 10]\nclass = a'),  
    Text(2863.6981132075475, 181.1999999999982, 'gini = 0.33\nsamples = 862\nvalue = [106  
7, 0, 267, 10]\nclass = a'),  
    Text(3032.1509433962265, 181.1999999999982, 'gini = 0.46\nsamples = 205\nvalue = [115,  
0, 206, 0]\nclass = c'),  
    Text(3790.1886792452833, 1268.4, 'OXY <= 4.675\ngini = 0.734\nsamples = 9639\nvalue =  
[3347, 2585, 4203, 5159]\nclass = d'),  
    Text(3453.283018867925, 906.0, 'OXY <= 1.415\ngini = 0.69\nsamples = 6912\nvalue = [87  
3, 2479, 2973, 4621]\nclass = d'),  
    Text(3284.8301886792456, 543.5999999999999, 'TOL <= 7.095\ngini = 0.635\nsamples = 1230  
\nvalue = [20, 963, 493, 491]\nclass = b'),  
    Text(3200.603773584906, 181.1999999999982, 'gini = 0.587\nsamples = 969\nvalue = [16,  
869, 281, 363]\nclass = b'),  
    Text(3369.0566037735853, 181.1999999999982, 'gini = 0.634\nsamples = 261\nvalue = [4,  
94, 212, 128]\nclass = c'),  
    Text(3621.735849056604, 543.5999999999999, 'TOL <= 11.005\ngini = 0.675\nsamples = 5682  
\nvalue = [853, 1516, 2480, 4130]\nclass = d'),  
    Text(3537.509433962264, 181.1999999999982, 'gini = 0.594\nsamples = 2795\nvalue = [18  
4, 927, 748, 2505]\nclass = d'),  
    Text(3705.9622641509436, 181.1999999999982, 'gini = 0.698\nsamples = 2887\nvalue = [66  
9, 589, 1732, 1625]\nclass = c'),  
    Text(4127.094339622642, 906.0, 'OXY <= 5.605\ngini = 0.58\nsamples = 2727\nvalue = [247  
4, 106, 1230, 538]\nclass = a'),  
    Text(3958.6415094339627, 543.5999999999999, 'NMHC <= 0.175\ngini = 0.669\nsamples = 695  
\nvalue = [437, 52, 414, 207]\nclass = a'),  
    Text(3874.415094339623, 181.1999999999982, 'gini = 0.526\nsamples = 295\nvalue = [292,  
5, 178, 23]\nclass = a'),  
    Text(4042.867924528302, 181.1999999999982, 'gini = 0.699\nsamples = 400\nvalue = [145,  
47, 236, 184]\nclass = c'),  
    Text(4295.5471698113215, 543.5999999999999, 'O_3 <= 9.305\ngini = 0.53\nsamples = 2032  
\nvalue = [2037, 54, 816, 331]\nclass = a'),  
    Text(4211.320754716981, 181.1999999999982, 'gini = 0.63\nsamples = 934\nvalue = [630,  
34, 618, 211]\nclass = a'),  
    Text(4379.773584905661, 181.1999999999982, 'gini = 0.332\nsamples = 1098\nvalue = [140  
7, 20, 198, 120]\nclass = a')]
```



# Conclusion

## Accuracy

Linear Regression:0.20446058578801884

Ridge Regression:0.20400914542939697

Lasso Regression:0.05954671177337789

ElasticNet Regression:0.10310757444971508

Logistic Regression:0.8480899292795158

Random Forest:0.7747286684902497

**Logistic Regression is suitable for this dataset**