

20104028

Heamnath N

Importing Libraries

In [1]:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

Importing Datasets

In [2]:

```
df=pd.read_csv("madrid_2003.csv")
df
```

Out[2]:

	date	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3	PM10	P
0	2003-03-01 01:00:00	NaN	1.72	NaN	NaN	NaN	73.900002	316.299988	NaN	10.550000	55.209999	NaN
1	2003-03-01 01:00:00	NaN	1.45	NaN	NaN	0.26	72.110001	250.000000	0.73	6.720000	52.389999	NaN
2	2003-03-01 01:00:00	NaN	1.57	NaN	NaN	NaN	80.559998	224.199997	NaN	21.049999	63.240002	NaN
3	2003-03-01 01:00:00	NaN	2.45	NaN	NaN	NaN	78.370003	450.399994	NaN	4.220000	67.839996	NaN
4	2003-03-01 01:00:00	NaN	3.26	NaN	NaN	NaN	96.250000	479.100006	NaN	8.460000	95.779999	NaN
...
243979	2003-10-01 00:00:00	0.20	0.16	2.01	3.17	0.02	31.799999	32.299999	1.68	34.049999	7.380000	1.2
243980	2003-10-01 00:00:00	0.32	0.08	0.36	0.72	NaN	10.450000	14.760000	1.00	34.610001	7.400000	0.1
243981	2003-10-01 00:00:00	NaN	NaN	NaN	NaN	0.07	34.639999	50.810001	NaN	32.160000	16.830000	NaN

	date	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3	PM10	PXY
243982	2003-10-01 00:00:00	NaN	NaN	NaN	NaN	0.07	32.580002	41.020000	NaN	NaN	13.570000	Na
243983	2003-10-01 00:00:00	1.00	0.29	2.15	6.41	0.07	37.150002	56.849998	2.28	21.480000	12.350000	2.4

243984 rows × 16 columns

Data Cleaning and Data Preprocessing

In [3]: `df=df.dropna()`In [4]: `df.columns`Out[4]: `Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3', 'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'], dtype='object')`In [5]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 33010 entries, 5 to 243983
Data columns (total 16 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   date        33010 non-null   object 
 1   BEN          33010 non-null   float64
 2   CO           33010 non-null   float64
 3   EBE          33010 non-null   float64
 4   MXY          33010 non-null   float64
 5   NMHC         33010 non-null   float64
 6   NO_2         33010 non-null   float64
 7   NOx          33010 non-null   float64
 8   OXY          33010 non-null   float64
 9   O_3           33010 non-null   float64
 10  PM10         33010 non-null   float64
 11  PXY          33010 non-null   float64
 12  SO_2         33010 non-null   float64
 13  TCH          33010 non-null   float64
 14  TOL          33010 non-null   float64
 15  station      33010 non-null   int64  
dtypes: float64(14), int64(1), object(1)
memory usage: 4.3+ MB
```

In [6]: `data=df[['CO', 'station']]`
`data`Out[6]:

	CO	station
5	1.94	28079006

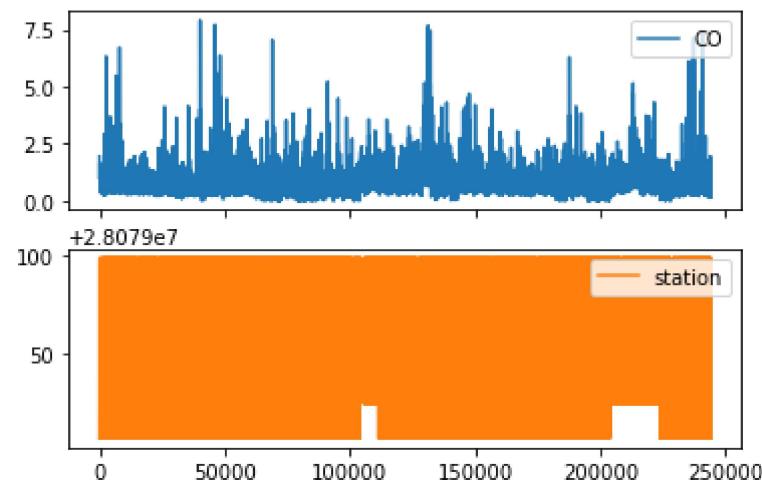
	CO	station
23	1.27	28079024
27	1.79	28079099
33	1.47	28079006
51	1.29	28079024
...
243955	0.41	28079099
243957	0.60	28079035
243961	0.82	28079006
243979	0.16	28079024
243983	0.29	28079099

33010 rows × 2 columns

Line chart

In [7]: `data.plot.line(subplots=True)`

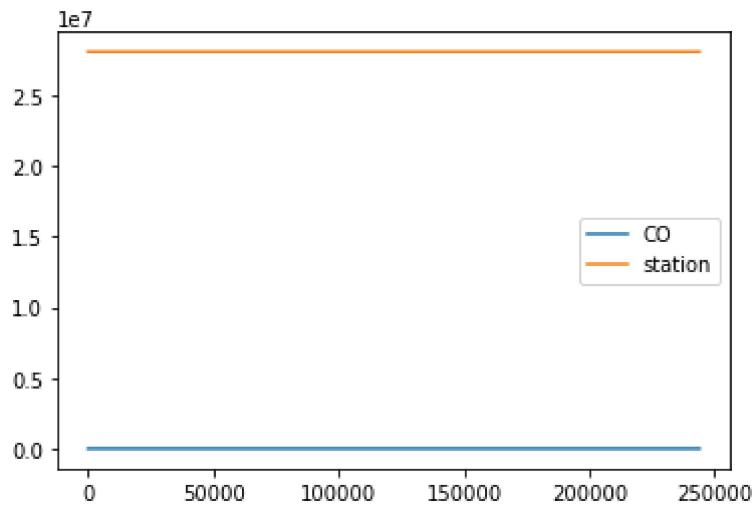
Out[7]: `array([<AxesSubplot:>, <AxesSubplot:>], dtype=object)`



Line chart

In [8]: `data.plot.line()`

Out[8]: `<AxesSubplot:>`

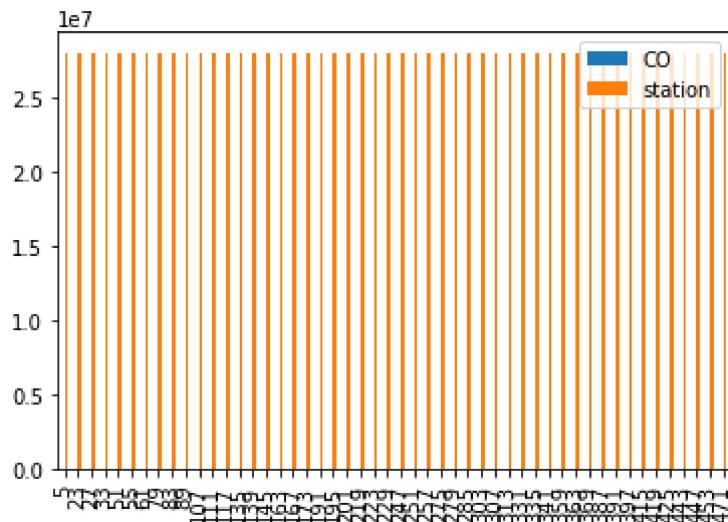


Bar chart

```
In [9]: b=data[0:50]
```

```
In [10]: b.plot.bar()
```

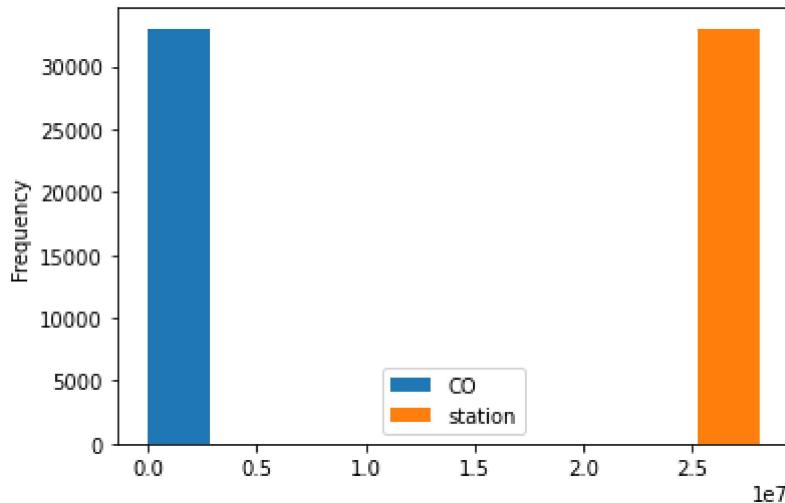
```
Out[10]: <AxesSubplot:>
```



Histogram

```
In [11]: data.plot.hist()
```

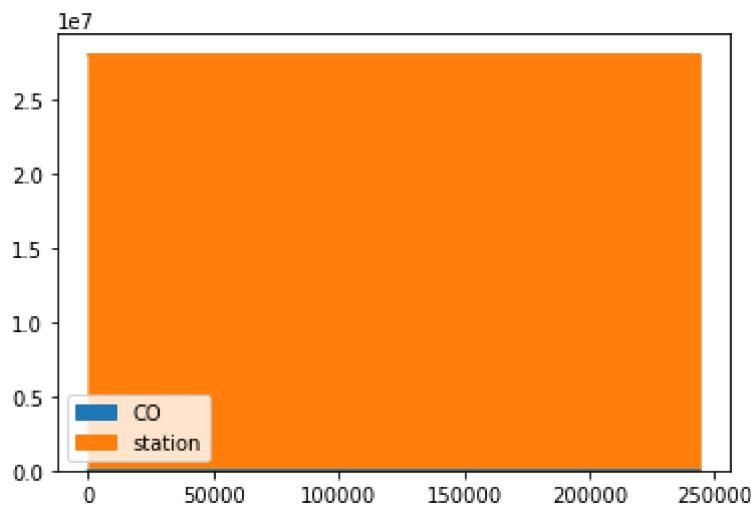
```
Out[11]: <AxesSubplot:ylabel='Frequency'>
```



Area chart

In [12]: `data.plot.area()`

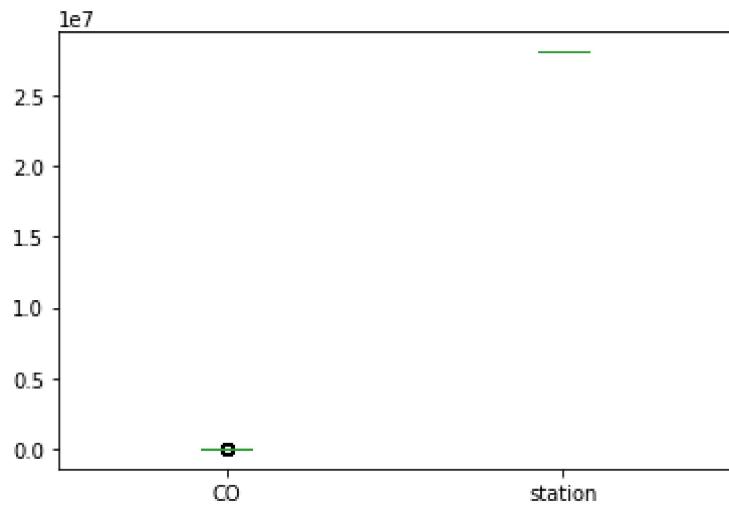
Out[12]: <AxesSubplot:>



Box chart

In [13]: `data.plot.box()`

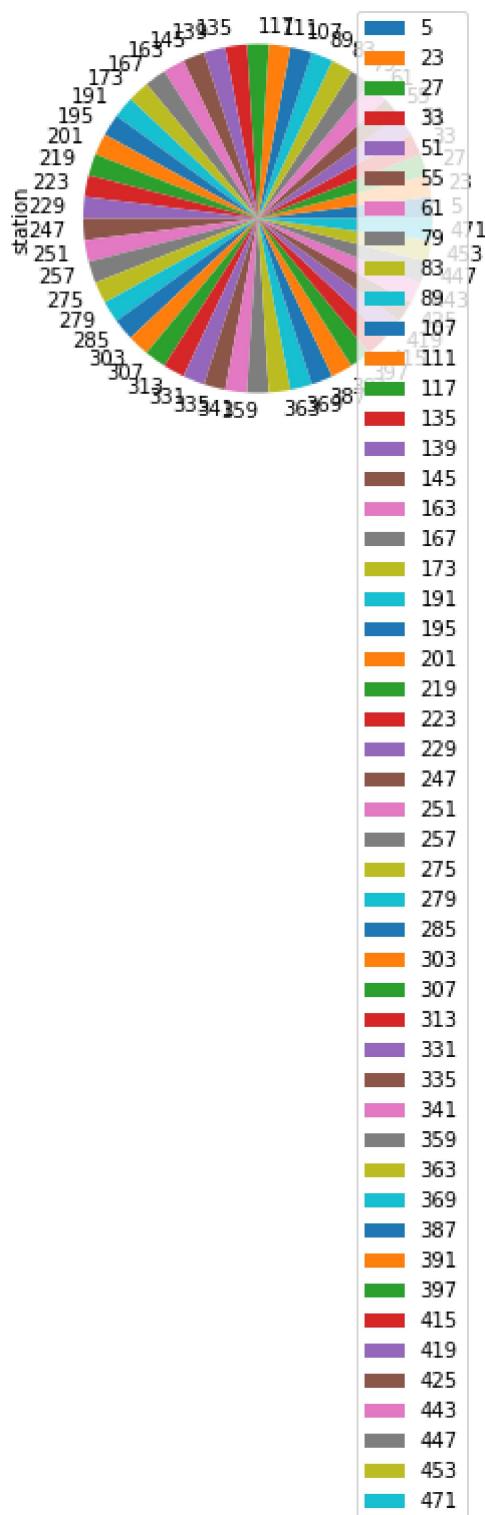
Out[13]: <AxesSubplot:>



Pie chart

```
In [14]: b.plot.pie(y='station' )
```

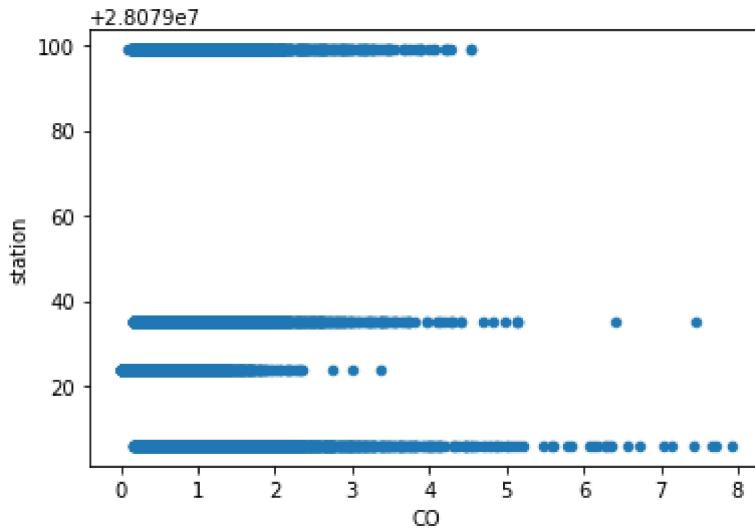
```
Out[14]: <AxesSubplot:ylabel='station'>
```



Scatter chart

In [15]: `data.plot.scatter(x='CO' ,y='station')`

Out[15]: <AxesSubplot:xlabel='CO', ylabel='station'>



In [16]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 33010 entries, 5 to 243983
Data columns (total 16 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   date        33010 non-null   object 
 1   BEN          33010 non-null   float64
 2   CO           33010 non-null   float64
 3   EBE          33010 non-null   float64
 4   MXY          33010 non-null   float64
 5   NMHC         33010 non-null   float64
 6   NO_2          33010 non-null   float64
 7   NOx          33010 non-null   float64
 8   OXY          33010 non-null   float64
 9   O_3           33010 non-null   float64
 10  PM10         33010 non-null   float64
 11  PXY          33010 non-null   float64
 12  SO_2          33010 non-null   float64
 13  TCH          33010 non-null   float64
 14  TOL          33010 non-null   float64
 15  station       33010 non-null   int64  
dtypes: float64(14), int64(1), object(1)
memory usage: 4.3+ MB
```

In [17]:

```
df.describe()
```

Out[17]:

	BEN	CO	EBE	MXY	NMHC	NO_2	NOx
count	33010.000000	33010.000000	33010.000000	33010.000000	33010.000000	33010.000000	33010.000000
mean	2.192633	0.759868	2.639726	5.838414	0.137177	57.328049	120.153671
std	2.064160	0.545999	2.825194	6.267296	0.127863	31.811082	104.521701
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.900000	0.430000	1.010000	1.880000	0.060000	34.529999	49.070000
50%	1.610000	0.620000	1.890000	4.070000	0.110000	55.105000	92.779991
75%	2.810000	0.930000	3.300000	7.530000	0.170000	76.160004	160.100000

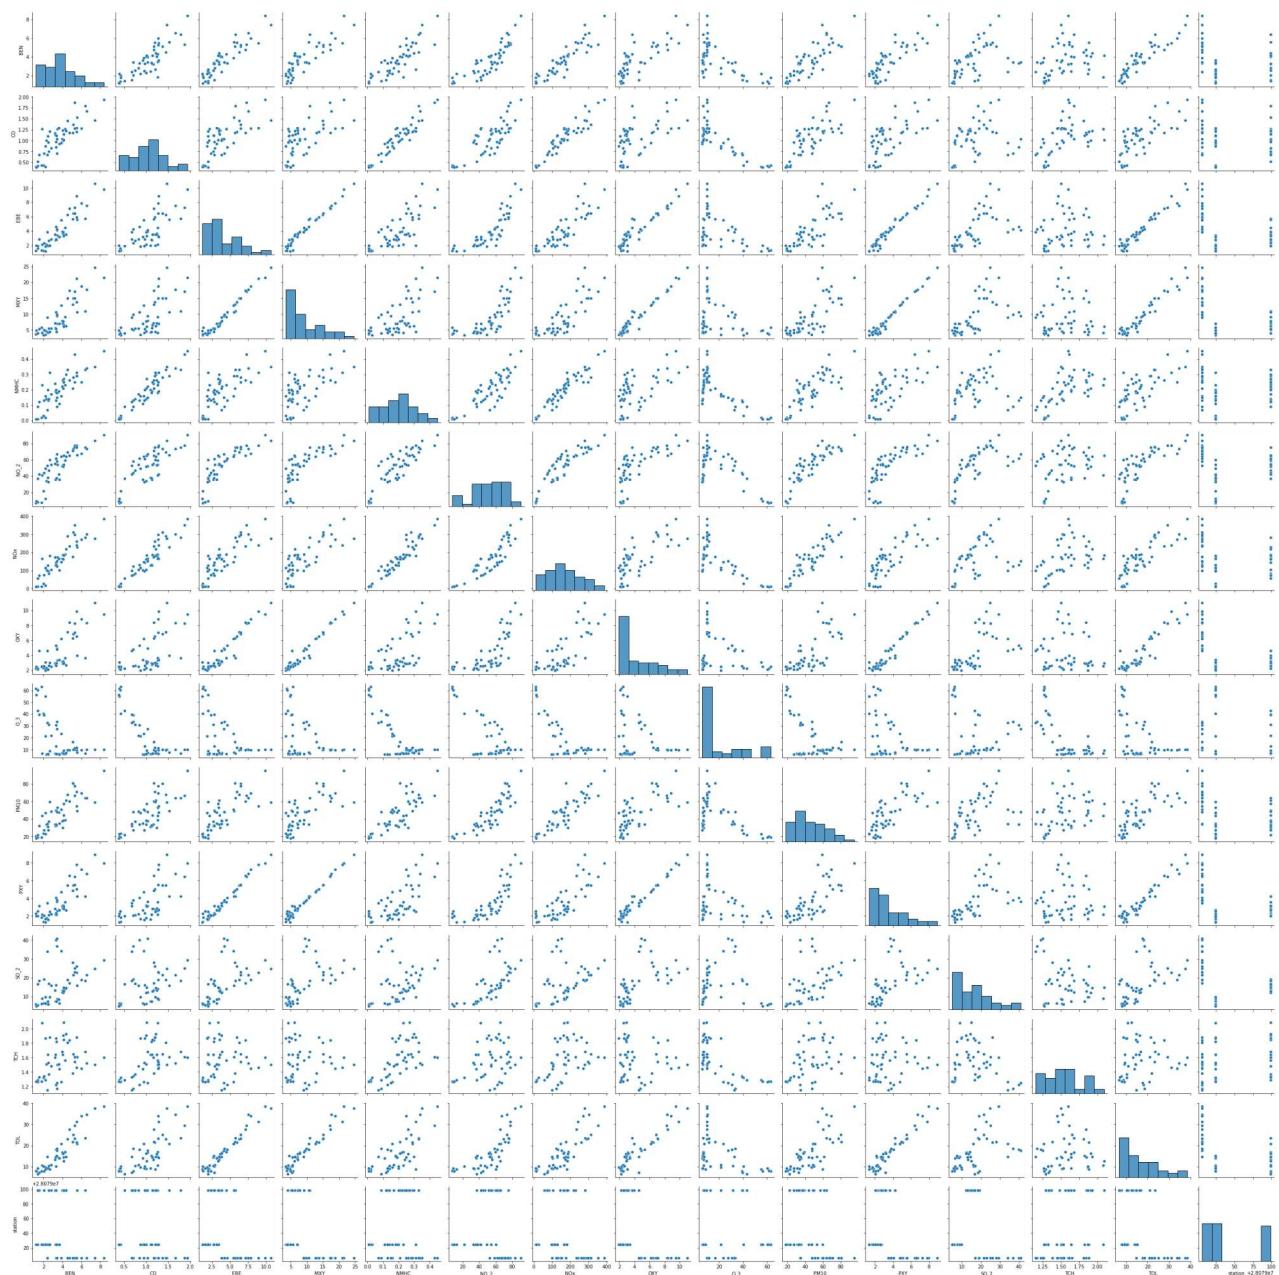
	BEN	CO	EBE	MXY	NMHC	NO_2	NOx
max	66.389999	7.920000	92.589996	177.600006	2.180000	342.700012	1246.000001

```
In [18]: df1=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3', 'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]
```

EDA AND VISUALIZATION

```
In [19]: sns.pairplot(df1[0:50])
```

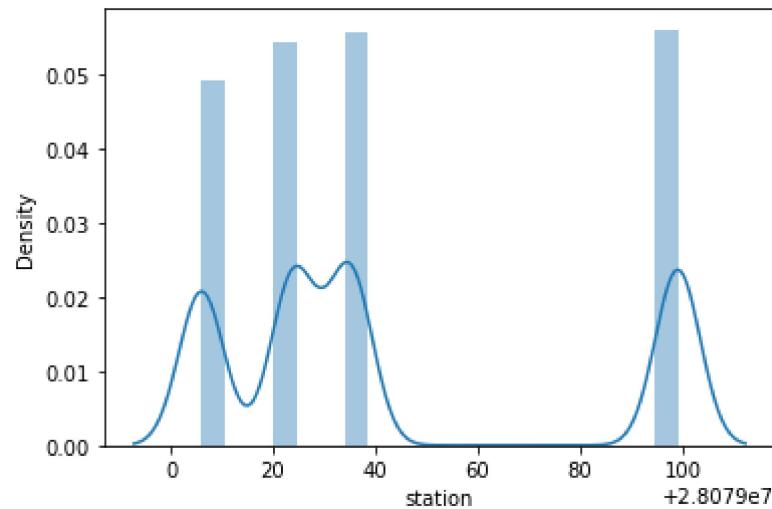
```
Out[19]: <seaborn.axisgrid.PairGrid at 0x19fe9c5f8b0>
```



```
In [20]: sns.distplot(df1['station'])
```

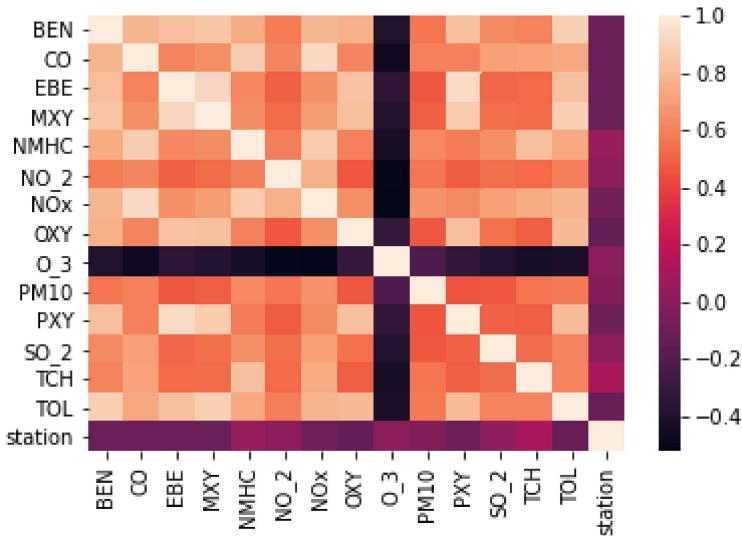
```
C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning:
`distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
    warnings.warn(msg, FutureWarning)
```

Out[20]: <AxesSubplot:xlabel='station', ylabel='Density'>



In [21]: `sns.heatmap(df1.corr())`

Out[21]: <AxesSubplot:>



TO TRAIN THE MODEL AND MODEL BUILDING

In [22]: `x=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3', 'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
y=df['station']`

In [23]: `from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)`

Linear Regression

```
In [24]: from sklearn.linear_model import LinearRegression  
lr=LinearRegression()  
lr.fit(x_train,y_train)
```

```
Out[24]: LinearRegression()
```

```
In [25]: lr.intercept_
```

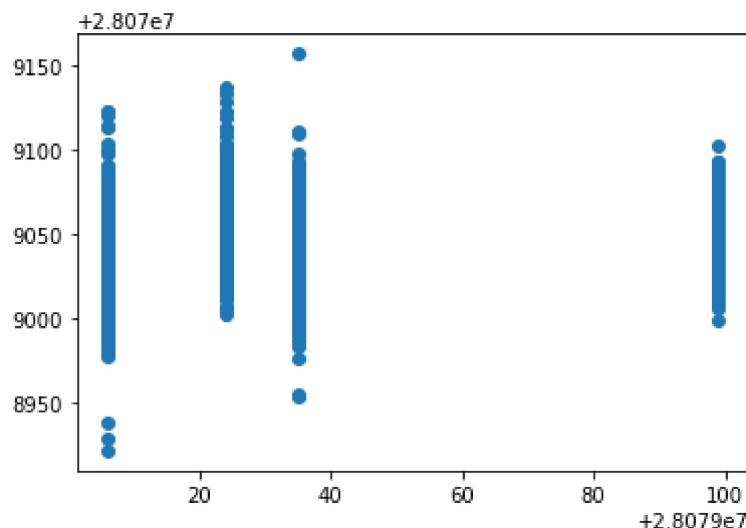
```
Out[25]: 28079002.996943984
```

```
In [26]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=[ 'Co-efficient'])  
coeff
```

	Co-efficient
BEN	1.276738
CO	-38.419943
EBE	-1.652659
MXY	0.120003
NMHC	155.388222
NO_2	0.155343
NOx	-0.067647
OXY	-1.316740
O_3	-0.013520
PM10	-0.057980
PXY	1.662017
SO_2	0.895247
TCH	34.063839
TOL	-0.836322

```
In [27]: prediction =lr.predict(x_test)  
plt.scatter(y_test,prediction)
```

```
Out[27]: <matplotlib.collections.PathCollection at 0x19ff9941790>
```



ACCURACY

```
In [28]: lr.score(x_test,y_test)
```

```
Out[28]: 0.16958450932868763
```

```
In [29]: lr.score(x_train,y_train)
```

```
Out[29]: 0.1787910184192012
```

Ridge and Lasso

```
In [30]: from sklearn.linear_model import Ridge,Lasso
```

```
In [31]: rr=Ridge(alpha=10)  
rr.fit(x_train,y_train)
```

```
Out[31]: Ridge(alpha=10)
```

Accuracy(Ridge)

```
In [32]: rr.score(x_test,y_test)
```

```
Out[32]: 0.16875553306978774
```

```
In [33]: rr.score(x_train,y_train)
```

```
Out[33]: 0.17770868147324226
```

```
In [34]: la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

Out[34]: Lasso(alpha=10)

```
In [35]: la.score(x_train,y_train)
```

Out[35]: 0.03709957832024513

Accuracy(Lasso)

```
In [36]: la.score(x_test,y_test)
```

Out[36]: 0.03394912089509239

```
In [37]: from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

Out[37]: ElasticNet()

```
In [38]: en.coef_
```

Out[38]: array([-0. , -0.32574759, 0.02119408, -0.02987462, 0.13340555,
 0.15031014, -0.0702183 , -1.30891183, -0.0421444 , 0.070547 ,
 0.23795802, 0.78604522, 1.58155984, -0.41258578])

```
In [39]: en.intercept_
```

Out[39]: 28079037.5312244

```
In [40]: prediction=en.predict(x_test)
```

```
In [41]: en.score(x_test,y_test)
```

Out[41]: 0.04461573097242899

Evaluation Metrics

```
In [42]: from sklearn import metrics
print(metrics.mean_absolute_error(y_test,prediction))
print(metrics.mean_squared_error(y_test,prediction))
print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
29.276741085597166  
1191.7744062400225  
34.52208577476197
```

Logistic Regression

```
In [43]: from sklearn.linear_model import LogisticRegression  
  
In [44]: feature_matrix=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOX', 'OXY', 'O_3',  
    'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]  
target_vector=df[ 'station']  
  
In [45]: feature_matrix.shape  
  
Out[45]: (33010, 14)  
  
In [46]: target_vector.shape  
  
Out[46]: (33010,)  
  
In [47]: from sklearn.preprocessing import StandardScaler  
  
In [48]: fs=StandardScaler().fit_transform(feature_matrix)  
  
In [49]: logr=LogisticRegression(max_iter=10000)  
logr.fit(fs,target_vector)  
  
Out[49]: LogisticRegression(max_iter=10000)  
  
In [50]: observation=[[1,2,3,4,5,6,7,8,9,10,11,12,13,14]]  
  
In [51]: prediction=logr.predict(observation)  
print(prediction)  
  
[28079035]  
  
In [52]: logr.classes_  
  
Out[52]: array([28079006, 28079024, 28079035, 28079099], dtype=int64)  
  
In [53]: logr.score(fs,target_vector)  
  
Out[53]: 0.7584974250227204
```

```
In [54]: logr.predict_proba(observation)[0][0]
```

```
Out[54]: 2.3306153265290618e-23
```

```
In [55]: logr.predict_proba(observation)
```

```
Out[55]: array([[2.33061533e-23, 1.44436075e-55, 1.00000000e+00, 6.68457491e-16]])
```

Random Forest

```
In [56]: from sklearn.ensemble import RandomForestClassifier
```

```
In [57]: rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[57]: RandomForestClassifier()
```

```
In [58]: parameters={'max_depth':[1,2,3,4,5],
                 'min_samples_leaf':[5,10,15,20,25],
                 'n_estimators':[10,20,30,40,50]
                }
```

```
In [59]: from sklearn.model_selection import GridSearchCV
grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")
grid_search.fit(x_train,y_train)
```

```
Out[59]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
                      param_grid={'max_depth': [1, 2, 3, 4, 5],
                                  'min_samples_leaf': [5, 10, 15, 20, 25],
                                  'n_estimators': [10, 20, 30, 40, 50]},
                      scoring='accuracy')
```

```
In [60]: grid_search.best_score_
```

```
Out[60]: 0.7247150023086772
```

```
In [61]: rfc_best=grid_search.best_estimator_
```

```
In [62]: from sklearn.tree import plot_tree
```

```
plt.figure(figsize=(80,40))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['a','b','c','d'])
```

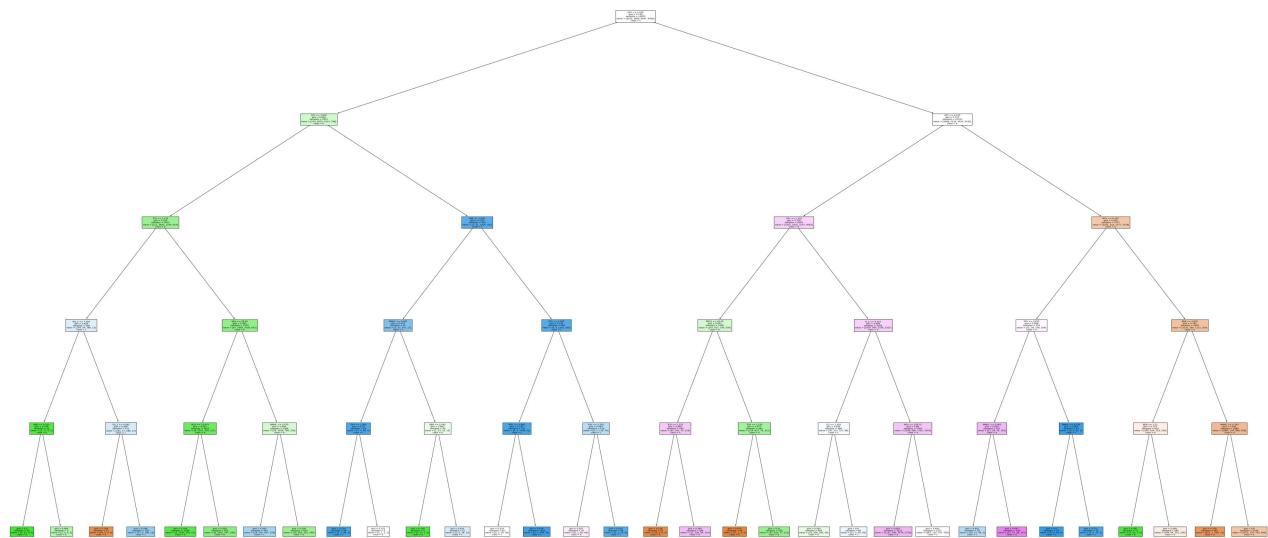
```
Out[62]: [Text(2232.0, 1993.2, 'OXY <= 1.105\nngini = 0.749\nnsamples = 14653\nvalue = [5312, 5809, 6047, 5939]\nnclass = c'),
          Text(1116.0, 1630.8000000000002, 'PXY <= 1.005\nngini = 0.604\nnsamples = 4551\nvalue = [214, 3675, 2513, 798]\nnclass = b')]
```

```

Text(558.0, 1268.4, 'TCH <= 1.235\ngini = 0.533\nsamples = 3622\nvalue = [211, 3643, 12
44, 616]\nclass = b'),
Text(279.0, 906.0, 'NO_2 <= 7.305\ngini = 0.629\nsamples = 246\nvalue = [144, 54, 186,
13]\nclass = c'),
Text(139.5, 543.5999999999999, 'BEN <= 0.525\ngini = 0.139\nsamples = 25\nvalue = [3, 3
7, 0, 0]\nclass = b'),
Text(69.75, 181.1999999999982, 'gini = 0.0\nsamples = 20\nvalue = [0, 32, 0, 0]\nclass
= b'),
Text(209.25, 181.1999999999982, 'gini = 0.469\nsamples = 5\nvalue = [3, 5, 0, 0]\nclas
s = b'),
Text(418.5, 543.5999999999999, 'SO_2 <= 4.995\ngini = 0.569\nsamples = 221\nvalue = [14
1, 17, 186, 13]\nclass = c'),
Text(348.75, 181.1999999999982, 'gini = 0.198\nsamples = 41\nvalue = [64, 8, 0, 0]\ncl
ass = a'),
Text(488.25, 181.1999999999982, 'gini = 0.498\nsamples = 180\nvalue = [77, 9, 186, 13]
\nclass = c'),
Text(837.0, 906.0, 'NOx <= 39.29\ngini = 0.492\nsamples = 3376\nvalue = [67, 3589, 105
8, 603]\nclass = b'),
Text(697.5, 543.5999999999999, 'NOx <= 22.675\ngini = 0.321\nsamples = 2010\nvalue =
[8, 2551, 350, 227]\nclass = b'),
Text(627.75, 181.1999999999982, 'gini = 0.164\nsamples = 1058\nvalue = [0, 1490, 103,
41]\nclass = b'),
Text(767.25, 181.1999999999982, 'gini = 0.459\nsamples = 952\nvalue = [8, 1061, 247, 1
86]\nclass = b'),
Text(976.5, 543.5999999999999, 'NMHC <= 0.075\ngini = 0.638\nsamples = 1366\nvalue = [5
9, 1038, 708, 376]\nclass = b'),
Text(906.75, 181.1999999999982, 'gini = 0.583\nsamples = 542\nvalue = [29, 209, 507, 1
28]\nclass = c'),
Text(1046.25, 181.1999999999982, 'gini = 0.538\nsamples = 824\nvalue = [30, 829, 201,
248]\nclass = b'),
Text(1674.0, 1268.4, 'EBE <= 1.005\ngini = 0.255\nsamples = 929\nvalue = [3, 32, 1269,
182]\nclass = c'),
Text(1395.0, 906.0, 'NMHC <= 0.085\ngini = 0.474\nsamples = 91\nvalue = [3, 27, 107, 1
7]\nclass = c'),
Text(1255.5, 543.5999999999999, 'OXY <= 1.065\ngini = 0.19\nsamples = 61\nvalue = [2,
4, 88, 4]\nclass = c'),
Text(1185.75, 181.1999999999982, 'gini = 0.124\nsamples = 56\nvalue = [0, 3, 86, 3]\ncl
ass = c'),
Text(1325.25, 181.1999999999982, 'gini = 0.722\nsamples = 5\nvalue = [2, 1, 2, 1]\ncla
ss = a'),
Text(1534.5, 543.5999999999999, 'BEN <= 0.545\ngini = 0.662\nsamples = 30\nvalue = [1,
23, 19, 13]\nclass = b'),
Text(1464.75, 181.1999999999982, 'gini = 0.105\nsamples = 9\nvalue = [0, 17, 1, 0]\ncl
ass = b'),
Text(1604.25, 181.1999999999982, 'gini = 0.633\nsamples = 21\nvalue = [1, 6, 18, 13]\n
class = c'),
Text(1953.0, 906.0, 'OXY <= 1.005\ngini = 0.224\nsamples = 838\nvalue = [0, 5, 1162, 16
5]\nclass = c'),
Text(1813.5, 543.5999999999999, 'MXY <= 2.685\ngini = 0.153\nsamples = 722\nvalue = [0,
4, 1048, 91]\nclass = c'),
Text(1743.75, 181.1999999999982, 'gini = 0.519\nsamples = 64\nvalue = [0, 2, 50, 50]\n
class = c'),
Text(1883.25, 181.1999999999982, 'gini = 0.079\nsamples = 658\nvalue = [0, 2, 998, 41]
\nclass = c'),
Text(2092.5, 543.5999999999999, 'PXY <= 1.755\ngini = 0.483\nsamples = 116\nvalue = [0,
1, 114, 74]\nclass = c'),
Text(2022.75, 181.1999999999982, 'gini = 0.507\nsamples = 76\nvalue = [0, 1, 60, 66]\n
class = d'),
Text(2162.25, 181.1999999999982, 'gini = 0.225\nsamples = 40\nvalue = [0, 0, 54, 8]\ncl
ass = c'),
Text(3348.0, 1630.800000000002, 'OXY <= 3.315\ngini = 0.725\nsamples = 10102\nvalue =
[5098, 2134, 3534, 5141]\nclass = d'),
Text(2790.0, 1268.4, 'PXY <= 1.125\ngini = 0.716\nsamples = 6365\nvalue = [1835, 1816,
2257, 4063]\nclass = d'),
Text(2511.0, 906.0, 'PM10 <= 20.39\ngini = 0.651\nsamples = 1080\nvalue = [207, 817, 14
]

```

```
9, 536]\nclass = b'),  
    Text(2371.5, 543.5999999999999, 'TCH <= 1.225\ngini = 0.663\nsamples = 439\nvalue = [9  
3, 190, 70, 324]\nclass = d'),  
    Text(2301.75, 181.1999999999982, 'gini = 0.136\nsamples = 53\nvalue = [76, 0, 6, 0]\nnc  
lass = a'),  
    Text(2441.25, 181.1999999999982, 'gini = 0.589\nsamples = 386\nvalue = [17, 190, 64, 3  
24]\nclass = d'),  
    Text(2650.5, 543.5999999999999, 'TCH <= 1.235\ngini = 0.571\nsamples = 641\nvalue = [11  
4, 627, 79, 212]\nclass = b'),  
    Text(2580.75, 181.1999999999982, 'gini = 0.073\nsamples = 48\nvalue = [77, 1, 2, 0]\nnc  
lass = a'),  
    Text(2720.25, 181.1999999999982, 'gini = 0.51\nsamples = 593\nvalue = [37, 626, 77, 21  
2]\nclass = b'),  
    Text(3069.0, 906.0, 'O_3 <= 6.315\ngini = 0.699\nsamples = 5285\nvalue = [1628, 999, 21  
08, 3527]\nclass = d'),  
    Text(2929.5, 543.5999999999999, 'CO <= 1.305\ngini = 0.635\nsamples = 486\nvalue = [30,  
313, 327, 98]\nclass = c'),  
    Text(2859.75, 181.1999999999982, 'gini = 0.582\nsamples = 400\nvalue = [19, 308, 262,  
40]\nclass = b'),  
    Text(2999.25, 181.1999999999982, 'gini = 0.6\nsamples = 86\nvalue = [11, 5, 65, 58]\nnc  
lass = c'),  
    Text(3208.5, 543.5999999999999, 'NOx <= 128.75\ngini = 0.68\nsamples = 4799\nvalue = [1  
598, 686, 1781, 3429]\nclass = d'),  
    Text(3138.75, 181.1999999999982, 'gini = 0.663\nsamples = 3522\nvalue = [1131, 566, 10  
74, 2729]\nclass = d'),  
    Text(3278.25, 181.1999999999982, 'gini = 0.693\nsamples = 1277\nvalue = [467, 120, 70  
7, 700]\nclass = c'),  
    Text(3906.0, 1268.4, 'MXY <= 6.165\ngini = 0.616\nsamples = 3737\nvalue = [3263, 318, 1  
277, 1078]\nclass = a'),  
    Text(3627.0, 906.0, 'PXY <= 4.675\ngini = 0.668\nsamples = 254\nvalue = [31, 58, 156, 1  
64]\nclass = d'),  
    Text(3487.5, 543.5999999999999, 'NMHC <= 0.085\ngini = 0.659\nsamples = 207\nvalue = [3  
1, 58, 74, 162]\nclass = d'),  
    Text(3417.75, 181.1999999999982, 'gini = 0.641\nsamples = 59\nvalue = [20, 21, 48, 5]  
\nclass = c'),  
    Text(3557.25, 181.1999999999982, 'gini = 0.497\nsamples = 148\nvalue = [11, 37, 26, 15  
7]\nclass = d'),  
    Text(3766.5, 543.5999999999999, 'NMHC <= 0.195\ngini = 0.046\nsamples = 47\nvalue = [0,  
0, 82, 2]\nclass = c'),  
    Text(3696.75, 181.1999999999982, 'gini = 0.0\nsamples = 35\nvalue = [0, 0, 63, 0]\nnc  
lass = c'),  
    Text(3836.25, 181.1999999999982, 'gini = 0.172\nsamples = 12\nvalue = [0, 0, 19, 2]\nnc  
lass = c'),  
    Text(4185.0, 906.0, 'BEN <= 2.625\ngini = 0.587\nsamples = 3483\nvalue = [3232, 260, 11  
21, 914]\nclass = a'),  
    Text(4045.5, 543.5999999999999, 'BEN <= 1.21\ngini = 0.717\nsamples = 603\nvalue = [34  
9, 106, 253, 256]\nclass = a'),  
    Text(3975.75, 181.1999999999982, 'gini = 0.087\nsamples = 23\nvalue = [0, 42, 2, 0]\nnc  
lass = b'),  
    Text(4115.25, 181.1999999999982, 'gini = 0.699\nsamples = 580\nvalue = [349, 64, 251,  
256]\nclass = a'),  
    Text(4324.5, 543.5999999999999, 'NMHC <= 0.145\ngini = 0.543\nsamples = 2880\nvalue =  
[2883, 154, 868, 658]\nclass = a'),  
    Text(4254.75, 181.1999999999982, 'gini = 0.278\nsamples = 681\nvalue = [911, 1, 165, 1  
3]\nclass = a'),  
    Text(4394.25, 181.1999999999982, 'gini = 0.6\nsamples = 2199\nvalue = [1972, 153, 703,  
645]\nclass = a')]
```



Conclusion

Accuracy

Linear Regression: 0.16958450932868763

Ridge Regression: 0.16875553306978774

Lasso Regression: 0.03394912089509239

ElasticNet Regression: 0.04461573097242899

Logistic Regression: 0.7584974250227204

Random Forest: 0.7247150023086772

Logistic Regression is suitable for this dataset