

20104028

Heamnath N

Importing Libraries

In [1]:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

Importing Datasets

In [2]:

```
df=pd.read_csv("madrid_2009.csv")
df
```

Out[2]:

	date	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3	PM10	PM
0	2009-10-01 01:00:00	NaN	0.27	NaN	NaN	NaN	39.889999	48.150002	NaN	50.680000	18.260000	↑
1	2009-10-01 01:00:00	NaN	0.22	NaN	NaN	NaN	21.230000	24.260000	NaN	55.880001	10.580000	↑
2	2009-10-01 01:00:00	NaN	0.18	NaN	NaN	NaN	31.230000	34.880001	NaN	49.060001	25.190001	↑
3	2009-10-01 01:00:00	0.95	0.33	1.43	2.68	0.25	55.180000	81.360001	1.57	36.669998	26.530001	€
4	2009-10-01 01:00:00	NaN	0.41	NaN	NaN	0.12	61.349998	76.260002	NaN	38.090000	23.760000	↑
...
215683	2009-06-01 00:00:00	0.50	0.22	0.39	0.75	0.09	22.000000	24.510000	1.00	82.239998	10.830000	↑
215684	2009-06-01 00:00:00	NaN	0.31	NaN	NaN	NaN	76.110001	101.099998	NaN	41.220001	9.920000	↑
215685	2009-06-01 00:00:00	0.13	NaN	0.86	NaN	0.23	81.050003	99.849998	NaN	24.830000	12.460000	€

	date	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3	PM10	PM25
215686	2009-06-01 00:00:00	0.21	NaN	2.96	NaN	0.10	72.419998	82.959999	NaN	NaN	13.030000	1
215687	2009-06-01 00:00:00	0.37	0.32	0.99	1.36	0.14	54.290001	64.480003	1.06	56.919998	15.360000	1

215688 rows × 17 columns

Data Cleaning and Data Preprocessing

In [3]: `df=df.dropna()`In [4]: `df.columns`Out[4]: `Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3', 'PM10', 'PM25', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'], dtype='object')`In [5]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 24717 entries, 3 to 215687
Data columns (total 17 columns):
 #   Column    Non-Null Count  Dtype  
--- 
 0   date      24717 non-null   object 
 1   BEN        24717 non-null   float64
 2   CO         24717 non-null   float64
 3   EBE        24717 non-null   float64
 4   MXY        24717 non-null   float64
 5   NMHC       24717 non-null   float64
 6   NO_2       24717 non-null   float64
 7   NOx        24717 non-null   float64
 8   OXY        24717 non-null   float64
 9   O_3         24717 non-null   float64
 10  PM10       24717 non-null   float64
 11  PM25       24717 non-null   float64
 12  PXY        24717 non-null   float64
 13  SO_2       24717 non-null   float64
 14  TCH        24717 non-null   float64
 15  TOL        24717 non-null   float64
 16  station    24717 non-null   int64  
dtypes: float64(15), int64(1), object(1)
memory usage: 3.4+ MB
```

In [6]: `data=df[['CO', 'station']]`
`data`Out[6]:

	CO	station
3	0.33	28079006

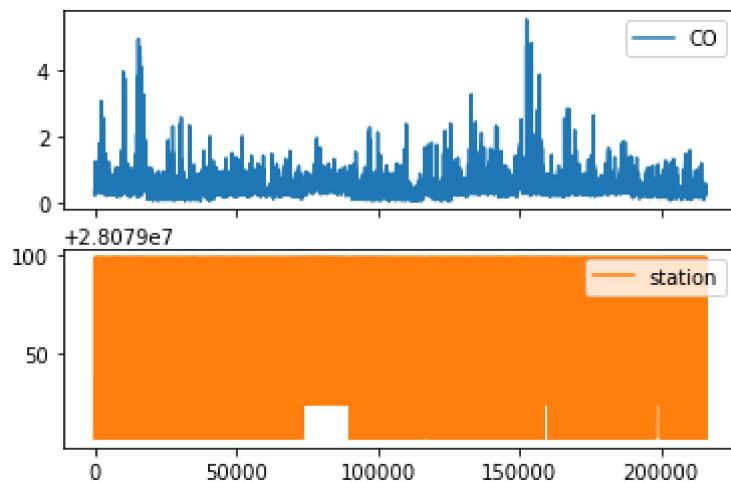
```
20 0.30 28079024
24 0.24 28079099
28 0.21 28079006
45 0.30 28079024
...
215659 0.27 28079024
215663 0.35 28079099
215667 0.29 28079006
215683 0.22 28079024
215687 0.32 28079099
```

24717 rows × 2 columns

Line chart

```
In [7]: data.plot.line(subplots=True)
```

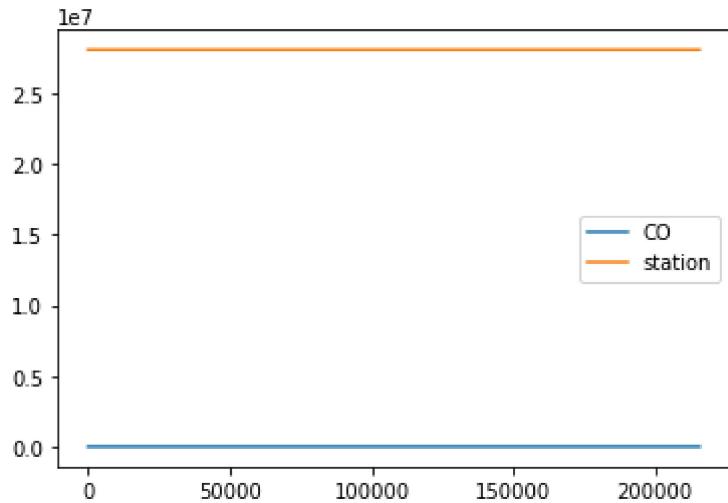
```
Out[7]: array([<AxesSubplot:>, <AxesSubplot:>], dtype=object)
```



Line chart

```
In [8]: data.plot.line()
```

```
Out[8]: <AxesSubplot:>
```

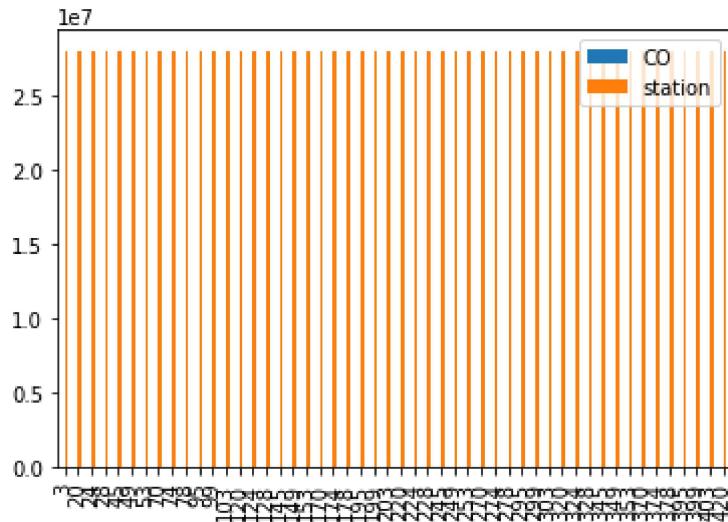


Bar chart

```
In [9]: b=data[0:50]
```

```
In [10]: b.plot.bar()
```

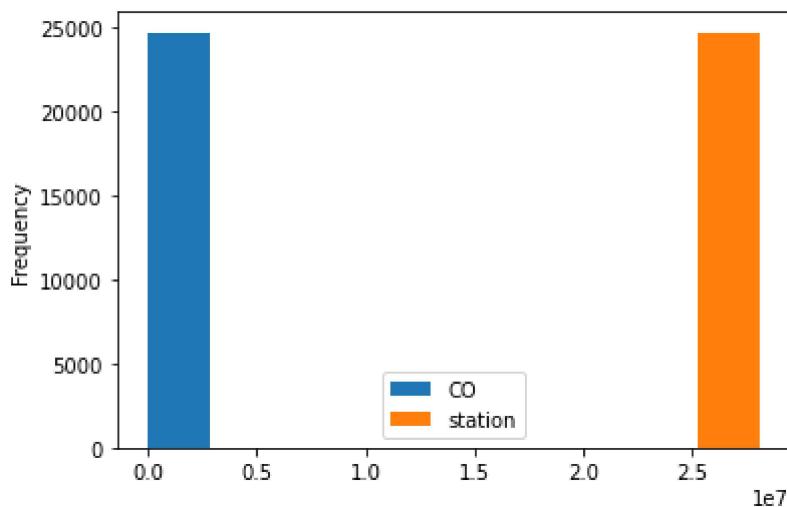
```
Out[10]: <AxesSubplot: >
```



Histogram

```
In [11]: data.plot.hist()
```

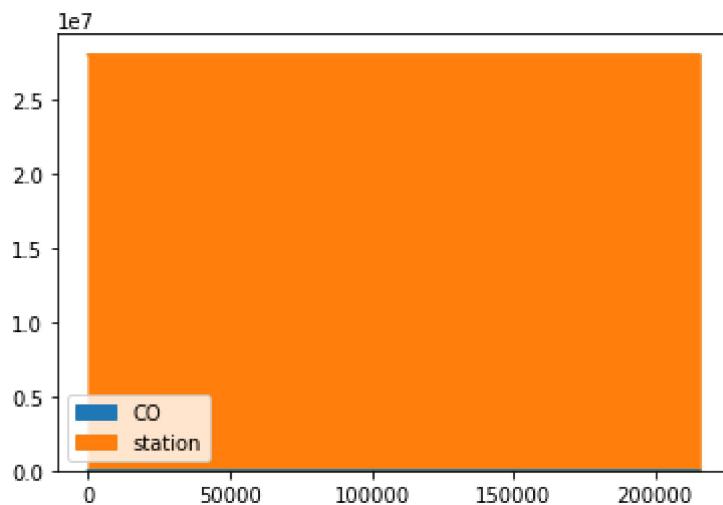
```
Out[11]: <AxesSubplot: ylabel='Frequency'>
```



Area chart

In [12]: `data.plot.area()`

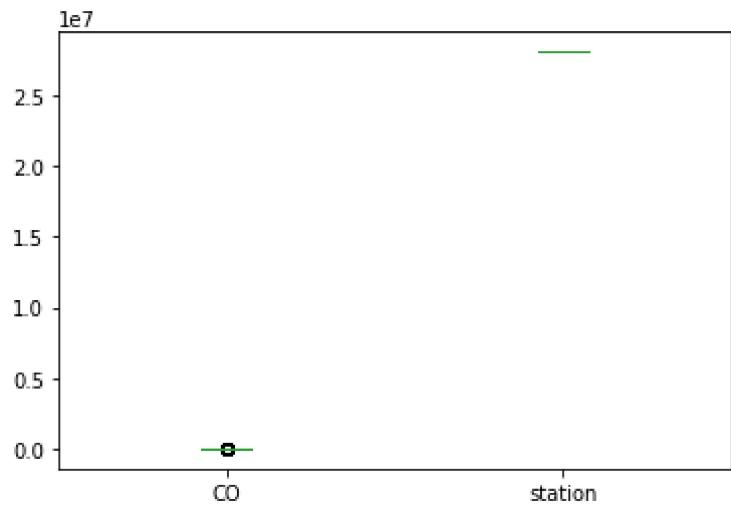
Out[12]: <AxesSubplot:>



Box chart

In [13]: `data.plot.box()`

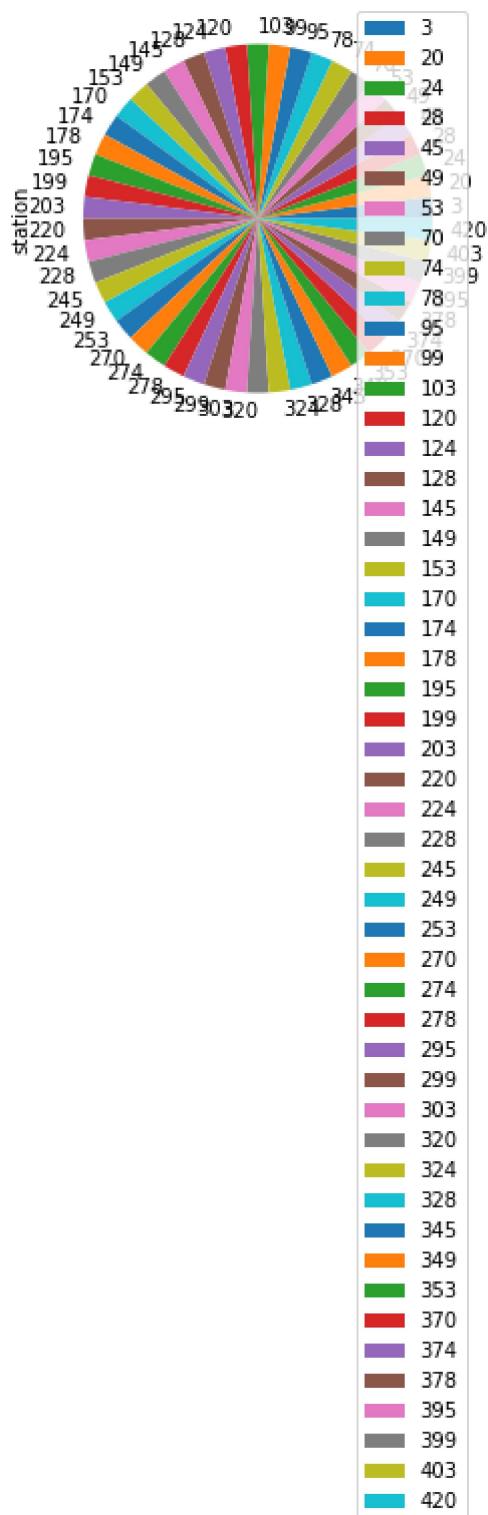
Out[13]: <AxesSubplot:>



Pie chart

```
In [14]: b.plot.pie(y='station' )
```

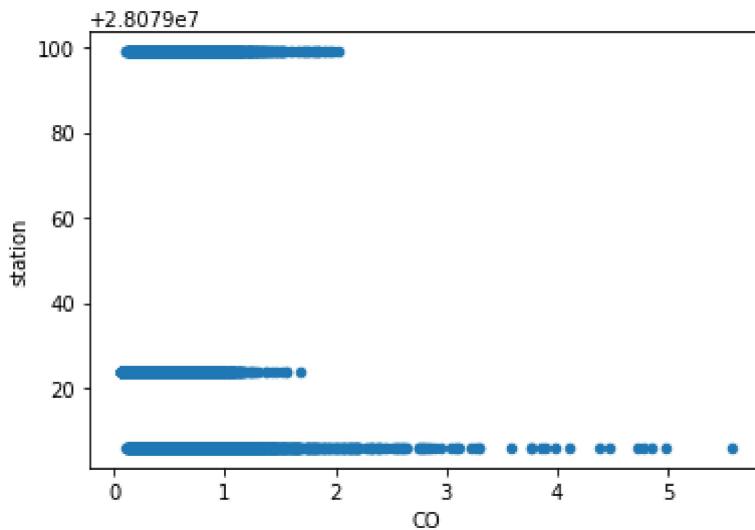
```
Out[14]: <AxesSubplot:ylabel='station'>
```



Scatter chart

```
In [15]: data.plot.scatter(x='CO' ,y='station')
```

```
Out[15]: <AxesSubplot:xlabel='CO', ylabel='station'>
```



In [16]:

`df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 24717 entries, 3 to 215687
Data columns (total 17 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   date        24717 non-null   object 
 1   BEN          24717 non-null   float64
 2   CO           24717 non-null   float64
 3   EBE          24717 non-null   float64
 4   MXY          24717 non-null   float64
 5   NMHC         24717 non-null   float64
 6   NO_2          24717 non-null   float64
 7   NOX          24717 non-null   float64
 8   OXY          24717 non-null   float64
 9   O_3           24717 non-null   float64
 10  PM10         24717 non-null   float64
 11  PM25         24717 non-null   float64
 12  PXY          24717 non-null   float64
 13  SO_2          24717 non-null   float64
 14  TCH           24717 non-null   float64
 15  TOL           24717 non-null   float64
 16  station       24717 non-null   int64  
dtypes: float64(15), int64(1), object(1)
memory usage: 3.4+ MB
```

In [17]:

`df.describe()`

Out[17]:

	BEN	CO	EBE	MXY	NMHC	NO_2	NO
count	24717.000000	24717.000000	24717.000000	24717.000000	24717.000000	24717.000000	24717.000000
mean	1.010583	0.448056	1.262430	2.244469	0.219582	55.563929	92.907181
std	1.007345	0.291706	1.074768	2.242214	0.141661	38.911677	91.985351
min	0.170000	0.060000	0.250000	0.240000	0.000000	0.600000	2.250000
25%	0.460000	0.270000	0.720000	0.990000	0.140000	26.510000	33.009991
50%	0.670000	0.370000	1.000000	1.490000	0.190000	47.930000	67.010000

	BEN	CO	EBE	MXY	NMHC	NO_2	NOx
75%	1.180000	0.570000	1.430000	2.820000	0.260000	76.269997	124.699997
max	22.379999	5.570000	47.669998	56.500000	2.580000	477.399994	1438.000000

In [18]:

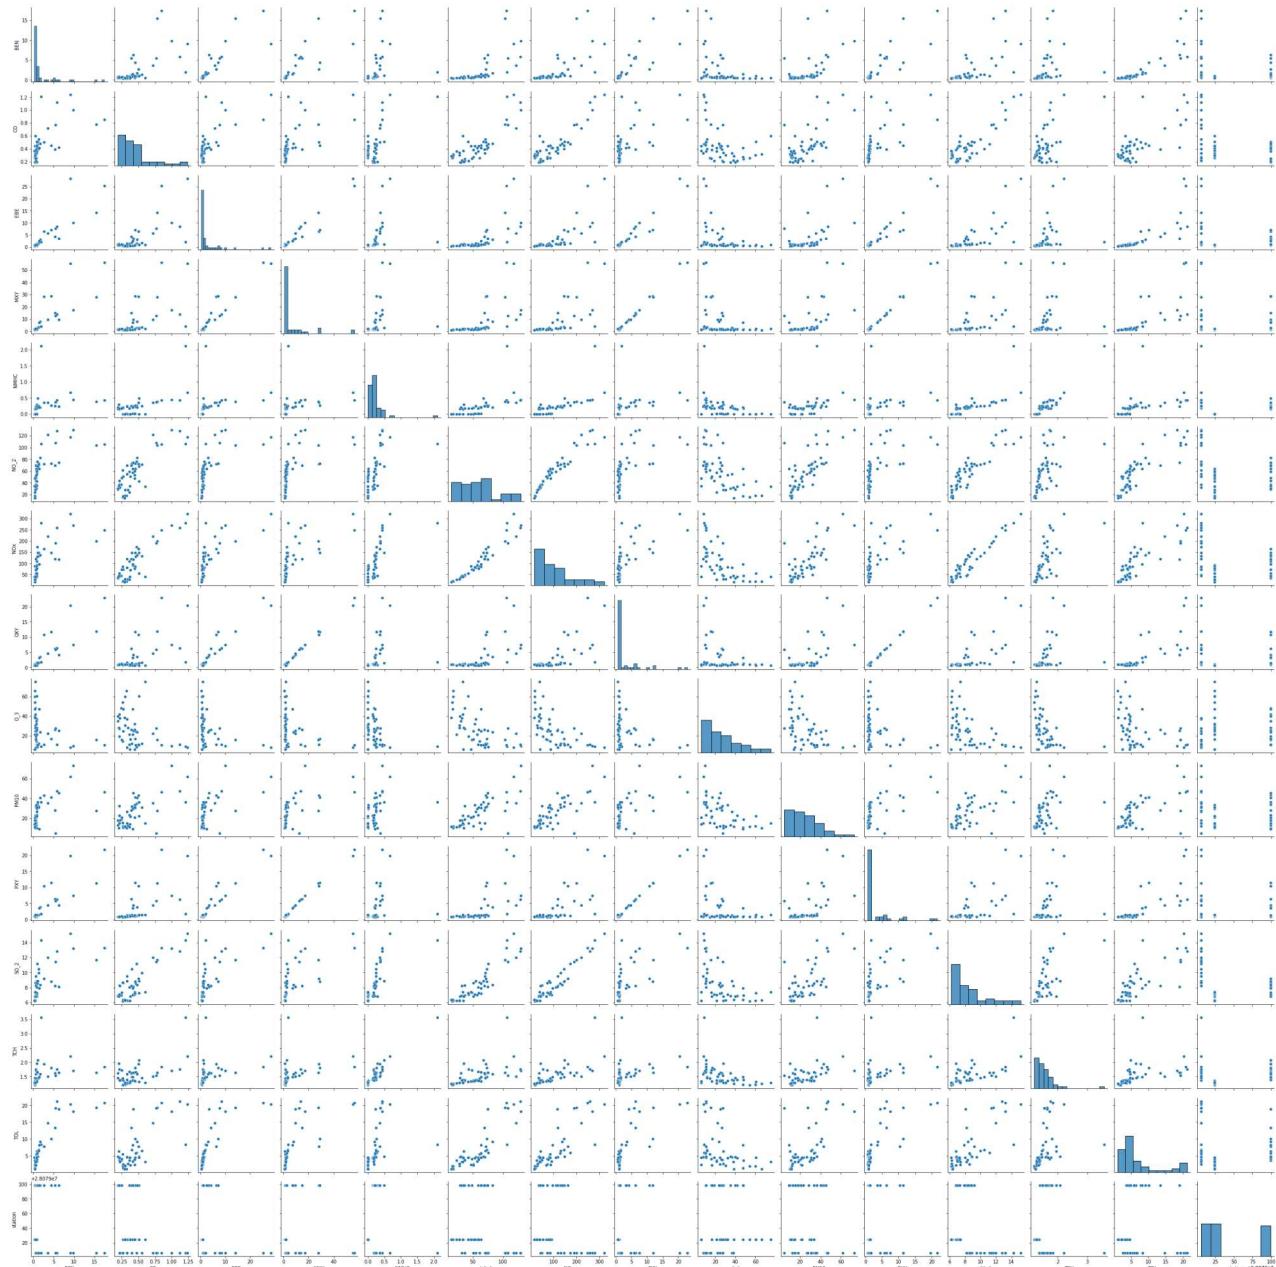
```
df1=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
       'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]
```

EDA AND VISUALIZATION

In [19]:

```
sns.pairplot(df1[0:50])
```

Out[19]: <seaborn.axisgrid.PairGrid at 0x1164aaf27c0>

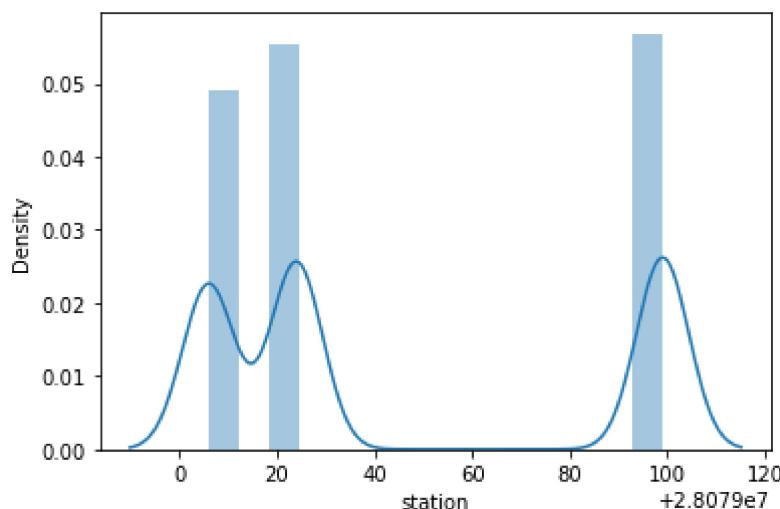


In [20]:

```
sns.distplot(df1['station'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
 warnings.warn(msg, FutureWarning)

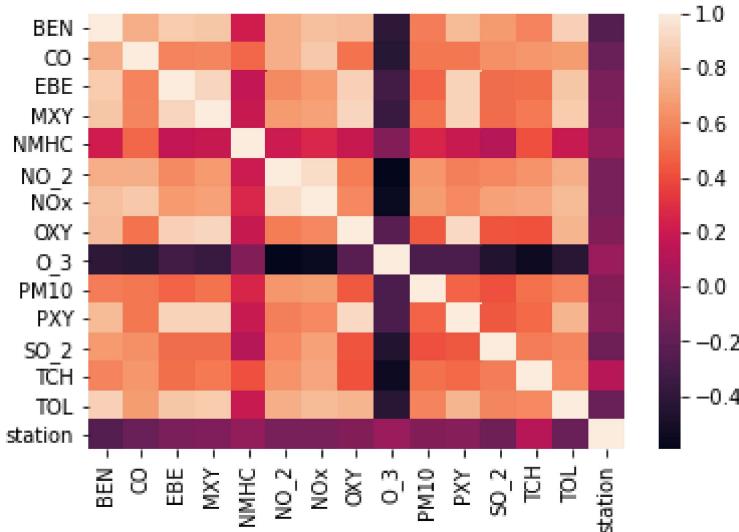
Out[20]: <AxesSubplot:xlabel='station', ylabel='Density'>



In [21]:

```
sns.heatmap(df1.corr())
```

Out[21]: <AxesSubplot:>



TO TRAIN THE MODEL AND MODEL BUILDING

In [22]:

```
x=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
       'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
y=df['station']
```

```
In [23]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

Linear Regression

```
In [24]: from sklearn.linear_model import LinearRegression
lr=LinearRegression()
lr.fit(x_train,y_train)
```

Out[24]: LinearRegression()

```
In [25]: lr.intercept_
```

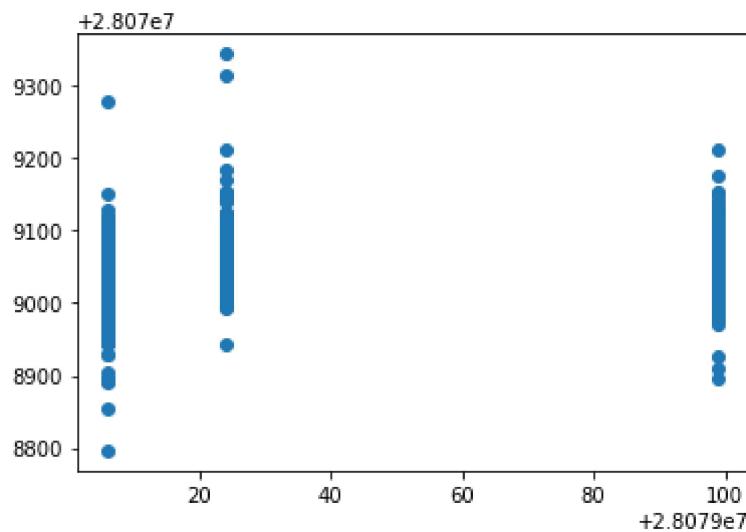
Out[25]: 28078892.665296998

```
In [26]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
coeff
```

	Co-efficient
BEN	-34.899220
CO	-28.765877
EBC	5.803456
MXY	-0.190398
NMHC	-19.084382
NO_2	-0.152268
NOx	0.177402
OXY	11.781123
O_3	0.025219
PM10	-0.050724
PXY	2.821465
SO_2	-0.387518
TCH	125.534611
TOL	-1.071650

```
In [27]: prediction =lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[27]: <matplotlib.collections.PathCollection at 0x11654cc98e0>



ACCURACY

```
In [28]: lr.score(x_test,y_test)
```

```
Out[28]: 0.27787771499038305
```

```
In [29]: lr.score(x_train,y_train)
```

```
Out[29]: 0.29072992168581624
```

Ridge and Lasso

```
In [30]: from sklearn.linear_model import Ridge,Lasso
```

```
In [31]: rr=Ridge(alpha=10)  
rr.fit(x_train,y_train)
```

```
Out[31]: Ridge(alpha=10)
```

Accuracy(Ridge)

```
In [32]: rr.score(x_test,y_test)
```

```
Out[32]: 0.2794352407863486
```

```
In [33]: rr.score(x_train,y_train)
```

```
Out[33]: 0.29038083663310943
```

```
In [34]: la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

Out[34]: Lasso(alpha=10)

```
In [35]: la.score(x_train,y_train)
```

Out[35]: 0.03875363015642119

Accuracy(Lasso)

```
In [36]: la.score(x_test,y_test)
```

Out[36]: 0.03495869696400822

```
In [37]: from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

Out[37]: ElasticNet()

```
In [38]: en.coef_
```

Out[38]: array([-7.00192409, -0.60387107, 0.1409905 , 2.08113437, -0.
-0.23004935, 0.13044212, 1.15080126, -0.14535201, 0.08614803,
 2.07352642, -0.84794817, 1.54538454, -1.9358594])

```
In [39]: en.intercept_
```

Out[39]: 28079064.34733883

```
In [40]: prediction=en.predict(x_test)
```

```
In [41]: en.score(x_test,y_test)
```

Out[41]: 0.10348190610556163

Evaluation Metrics

```
In [42]: from sklearn import metrics
print(metrics.mean_absolute_error(y_test,prediction))
print(metrics.mean_squared_error(y_test,prediction))
print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
35.86803561911397  
1462.6597580464495  
38.244735036949194
```

Logistic Regression

```
In [43]: from sklearn.linear_model import LogisticRegression
```

```
In [44]: feature_matrix=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOX', 'OXY', 'O_3',  
    'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]  
target_vector=df[ 'station']
```

```
In [45]: feature_matrix.shape
```

```
Out[45]: (24717, 14)
```

```
In [46]: target_vector.shape
```

```
Out[46]: (24717,)
```

```
In [47]: from sklearn.preprocessing import StandardScaler
```

```
In [48]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [49]: logr=LogisticRegression(max_iter=10000)  
logr.fit(fs,target_vector)
```

```
Out[49]: LogisticRegression(max_iter=10000)
```

```
In [50]: observation=[[1,2,3,4,5,6,7,8,9,10,11,12,13,14]]
```

```
In [51]: prediction=logr.predict(observation)  
print(prediction)
```

```
[28079099]
```

```
In [52]: logr.classes_
```

```
Out[52]: array([28079006, 28079024, 28079099], dtype=int64)
```

```
In [53]: logr.score(fs,target_vector)
```

```
Out[53]: 0.8951733624630821
```

```
In [54]: logr.predict_proba(observation)[0][0]
```

```
Out[54]: 5.447205522232353e-13
```

```
In [55]: logr.predict_proba(observation)
```

```
Out[55]: array([[5.44720552e-13, 8.28692830e-44, 1.00000000e+00]])
```

Random Forest

```
In [56]: from sklearn.ensemble import RandomForestClassifier
```

```
In [57]: rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[57]: RandomForestClassifier()
```

```
In [58]: parameters={'max_depth':[1,2,3,4,5],
                 'min_samples_leaf':[5,10,15,20,25],
                 'n_estimators':[10,20,30,40,50]
                }
```

```
In [59]: from sklearn.model_selection import GridSearchCV
grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")
grid_search.fit(x_train,y_train)
```

```
Out[59]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
                      param_grid={'max_depth': [1, 2, 3, 4, 5],
                                  'min_samples_leaf': [5, 10, 15, 20, 25],
                                  'n_estimators': [10, 20, 30, 40, 50]},
                      scoring='accuracy')
```

```
In [60]: grid_search.best_score_
```

```
Out[60]: 0.9047453433496612
```

```
In [61]: rfc_best=grid_search.best_estimator_
```

```
In [62]: from sklearn.tree import plot_tree
```

```
plt.figure(figsize=(80,40))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['a','b','c','d'])
```

```
Out[62]: [Text(2315.700000000003, 1993.2, 'BEN <= 1.355\ngini = 0.665\nsamples = 10945\nvalue =
[5162, 5955, 6184]\nklass = c'),
Text(1190.4, 1630.800000000002, 'SO_2 <= 6.855\ngini = 0.63\nsamples = 8688\nvalue =
[2442, 5747, 5585]\nklass = b')]
```

```

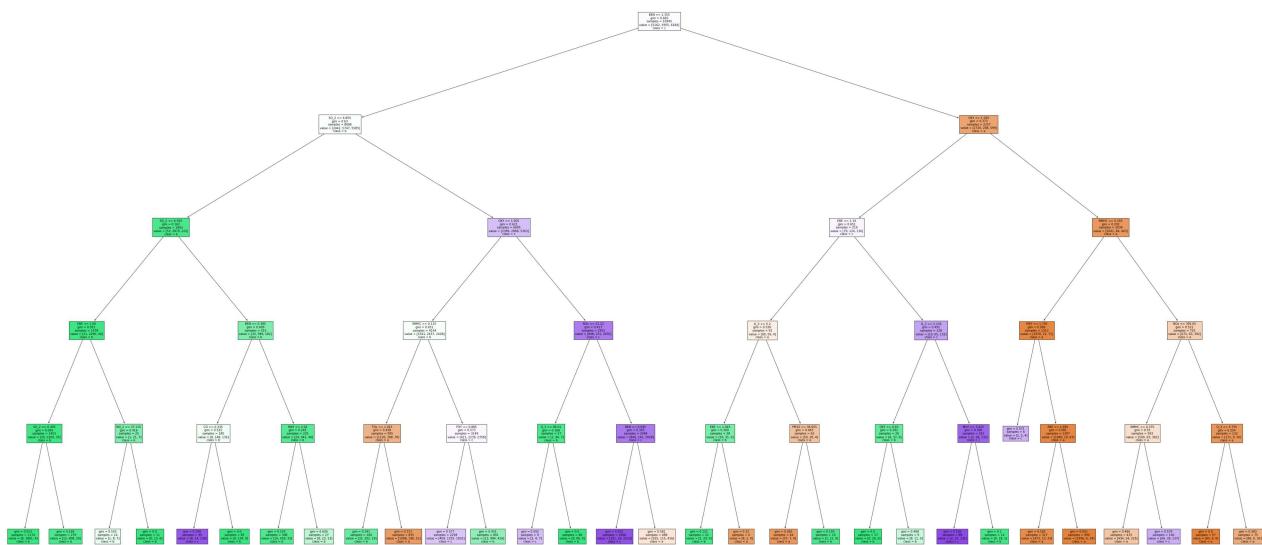
Text(595.2, 1268.4, 'SO_2 <= 6.565\ngini = 0.162\nsamples = 1993\nvalue = [53, 2879, 22
2]\nnclass = b'),
Text(297.6, 906.0, 'EBE <= 1.06\ngini = 0.051\nsamples = 1478\nvalue = [21, 2290, 40]\n
class = b'),
Text(148.8, 543.599999999999, 'SO_2 <= 6.405\ngini = 0.045\nsamples = 1453\nvalue = [2
0, 2269, 33]\nnclass = b'),
Text(74.4, 181.1999999999982, 'gini = 0.013\nsamples = 1174\nvalue = [8, 1861, 4]\ncla
ss = b'),
Text(223.2000000000002, 181.1999999999982, 'gini = 0.169\nsamples = 279\nvalue = [12,
408, 29]\nnclass = b'),
Text(446.4000000000003, 543.599999999999, 'NO_2 <= 37.115\ngini = 0.416\nsamples = 25
\nvalue = [1, 21, 7]\nnclass = b'),
Text(372.0, 181.1999999999982, 'gini = 0.555\nsamples = 14\nvalue = [1, 8, 7]\nnclass =
b'),
Text(520.800000000001, 181.1999999999982, 'gini = 0.0\nsamples = 11\nvalue = [0, 13,
0]\nnclass = b'),
Text(892.800000000001, 906.0, 'BEN <= 0.385\ngini = 0.409\nsamples = 515\nvalue = [32,
589, 182]\nnclass = b'),
Text(744.0, 543.599999999999, 'CO <= 0.335\ngini = 0.525\nsamples = 180\nvalue = [8, 1
48, 136]\nnclass = b'),
Text(669.6, 181.1999999999982, 'gini = 0.249\nsamples = 95\nvalue = [8, 14, 136]\nclas
s = c'),
Text(818.400000000001, 181.1999999999982, 'gini = 0.0\nsamples = 85\nvalue = [0, 134,
0]\nnclass = b'),
Text(1041.600000000001, 543.599999999999, 'MXY <= 2.02\ngini = 0.245\nsamples = 335\n
value = [24, 441, 46]\nnclass = b'),
Text(967.2, 181.1999999999982, 'gini = 0.193\nsamples = 308\nvalue = [16, 418, 33]\ncl
ass = b'),
Text(1116.0, 181.1999999999982, 'gini = 0.606\nsamples = 27\nvalue = [8, 23, 13]\nclas
s = b'),
Text(1785.600000000001, 1268.4, 'OXY <= 1.005\ngini = 0.621\nsamples = 6695\nvalue =
[2389, 2868, 5363]\nnclass = c'),
Text(1488.0, 906.0, 'NMHC <= 0.115\ngini = 0.651\nsamples = 4144\nvalue = [1541, 2637,
2428]\nnclass = b'),
Text(1339.2, 543.599999999999, 'TOL <= 1.015\ngini = 0.438\nsamples = 995\nvalue = [11
20, 398, 70]\nnclass = a'),
Text(1264.800000000002, 181.1999999999982, 'gini = 0.341\nsamples = 160\nvalue = [32,
202, 19]\nnclass = b'),
Text(1413.600000000001, 181.1999999999982, 'gini = 0.313\nsamples = 835\nvalue = [108
8, 196, 51]\nnclass = a'),
Text(1636.800000000002, 543.599999999999, 'PXY <= 0.885\ngini = 0.573\nsamples = 3149
\nvalue = [421, 2239, 2358]\nnclass = c'),
Text(1562.4, 181.1999999999982, 'gini = 0.577\nsamples = 2258\nvalue = [409, 1255, 193
2]\nnclass = c'),
Text(1711.2, 181.1999999999982, 'gini = 0.431\nsamples = 891\nvalue = [12, 984, 426]\n
class = b'),
Text(2083.200000000003, 906.0, 'NOx <= 21.22\ngini = 0.417\nsamples = 2551\nvalue = [8
48, 231, 2935]\nnclass = c'),
Text(1934.4, 543.599999999999, 'O_3 <= 86.01\ngini = 0.168\nsamples = 57\nvalue = [2,
90, 7]\nnclass = b'),
Text(1860.000000000002, 181.1999999999982, 'gini = 0.592\nsamples = 9\nvalue = [2, 4,
7]\nnclass = c'),
Text(2008.800000000002, 181.1999999999982, 'gini = 0.0\nsamples = 48\nvalue = [0, 86,
0]\nnclass = b'),
Text(2232.0, 543.599999999999, 'BEN <= 0.945\ngini = 0.393\nsamples = 2494\nvalue = [8
46, 141, 2928]\nnclass = c'),
Text(2157.600000000004, 181.1999999999982, 'gini = 0.201\nsamples = 1806\nvalue = [29
1, 26, 2512]\nnclass = c'),
Text(2306.4, 181.1999999999982, 'gini = 0.581\nsamples = 688\nvalue = [555, 115, 416]
\nnclass = a'),
Text(3441.000000000005, 1630.800000000002, 'OXY <= 1.005\ngini = 0.373\nsamples = 225
7\nvalue = [2720, 208, 599]\nnclass = a'),
Text(2976.0, 1268.4, 'EBE <= 1.14\ngini = 0.651\nsamples = 219\nvalue = [79, 124, 136]
\nnclass = c'),
Text(2678.4, 906.0, 'O_3 <= 5.2\ngini = 0.526\nsamples = 91\nvalue = [69, 59, 4]\nnclass

```

```

= a'),
Text(2529.600000000004, 543.599999999999, 'EBE <= 1.045\ngini = 0.369\nsamples = 28\n
value = [10, 31, 0]\nclass = b'),
Text(2455.200000000003, 181.1999999999982, 'gini = 0.121\nsamples = 22\nvalue = [2, 2
9, 0]\nclass = b'),
Text(2604.0, 181.1999999999982, 'gini = 0.32\nsamples = 6\nvalue = [8, 2, 0]\nclass =
a'),
Text(2827.200000000003, 543.599999999999, 'PM10 <= 38.055\ngini = 0.483\nsamples = 63
\nvalue = [59, 28, 4]\nclass = a'),
Text(2752.8, 181.1999999999982, 'gini = 0.283\nsamples = 44\nvalue = [57, 7, 4]\nclass
= a'),
Text(2901.600000000004, 181.1999999999982, 'gini = 0.159\nsamples = 19\nvalue = [2, 2
1, 0]\nclass = b'),
Text(3273.600000000004, 906.0, 'O_3 <= 5.235\ngini = 0.492\nsamples = 128\nvalue = [1
0, 65, 132]\nclass = c'),
Text(3124.8, 543.599999999999, 'OXY <= 0.63\ngini = 0.292\nsamples = 26\nvalue = [8, 3
7, 0]\nclass = b'),
Text(3050.4, 181.1999999999982, 'gini = 0.0\nsamples = 17\nvalue = [0, 26, 0]\nclass =
b'),
Text(3199.200000000003, 181.1999999999982, 'gini = 0.488\nsamples = 9\nvalue = [8, 1
1, 0]\nclass = b'),
Text(3422.4, 543.599999999999, 'MXY <= 3.025\ngini = 0.306\nsamples = 102\nvalue = [2,
28, 132]\nclass = c'),
Text(3348.000000000005, 181.1999999999982, 'gini = 0.156\nsamples = 88\nvalue = [2, 1
0, 131]\nclass = c'),
Text(3496.8, 181.1999999999982, 'gini = 0.1\nsamples = 14\nvalue = [0, 18, 1]\nclass =
b'),
Text(3906.000000000005, 1268.4, 'NMHC <= 0.265\ngini = 0.292\nsamples = 2038\nvalue =
[2641, 84, 463]\nclass = a'),
Text(3645.600000000004, 906.0, 'MXY <= 1.745\ngini = 0.086\nsamples = 1313\nvalue = [1
970, 21, 71]\nclass = a'),
Text(3571.200000000003, 543.599999999999, 'gini = 0.571\nsamples = 6\nvalue = [1, 2,
4]\nclass = c'),
Text(3720.000000000005, 543.599999999999, 'EBE <= 1.695\ngini = 0.081\nsamples = 1307
\nvalue = [1969, 19, 67]\nclass = a'),
Text(3645.600000000004, 181.1999999999982, 'gini = 0.165\nsamples = 317\nvalue = [47
3, 13, 33]\nclass = a'),
Text(3794.4, 181.1999999999982, 'gini = 0.051\nsamples = 990\nvalue = [1496, 6, 34]\n
lass = a'),
Text(4166.400000000001, 906.0, 'NOx <= 395.65\ngini = 0.521\nsamples = 725\nvalue = [67
1, 63, 392]\nclass = a'),
Text(4017.600000000004, 543.599999999999, 'NMHC <= 0.375\ngini = 0.55\nsamples = 593
\nvalue = [500, 63, 362]\nclass = a'),
Text(3943.200000000003, 181.1999999999982, 'gini = 0.486\nsamples = 433\nvalue = [43
4, 24, 225]\nclass = a'),
Text(4092.000000000005, 181.1999999999982, 'gini = 0.579\nsamples = 160\nvalue = [66,
39, 137]\nclass = c'),
Text(4315.200000000001, 543.599999999999, 'O_3 <= 5.755\ngini = 0.254\nsamples = 132\n
value = [171, 0, 30]\nclass = a'),
Text(4240.8, 181.1999999999982, 'gini = 0.0\nsamples = 57\nvalue = [85, 0, 0]\nclass =
a'),
Text(4389.6, 181.1999999999982, 'gini = 0.383\nsamples = 75\nvalue = [86, 0, 30]\nclas
s = a')]

```



Conclusion

Accuracy

In [63]:

```
print("Linear Regression:",lr.score(x_test,y_test))
print("Ridge Regression:",rr.score(x_test,y_test))
print("Lasso Regression",la.score(x_test,y_test))
print("ElasticNet Regression:",en.score(x_test,y_test))
print("Logistic Regression:",logr.score(fs,target_vector))
print("Random Forest:",grid_search.best_score_)
```

```
Linear Regression: 0.27787771499038305
Ridge Regression: 0.2794352407863486
Lasso Regression 0.03495869696400822
ElasticNet Regression: 0.10348190610556163
Logistic Regression: 0.8951733624630821
Random Forest: 0.9047453433496612
```

Random Forest is suitable for this dataset